



Image Analysis

Rasmus R. Paulsen

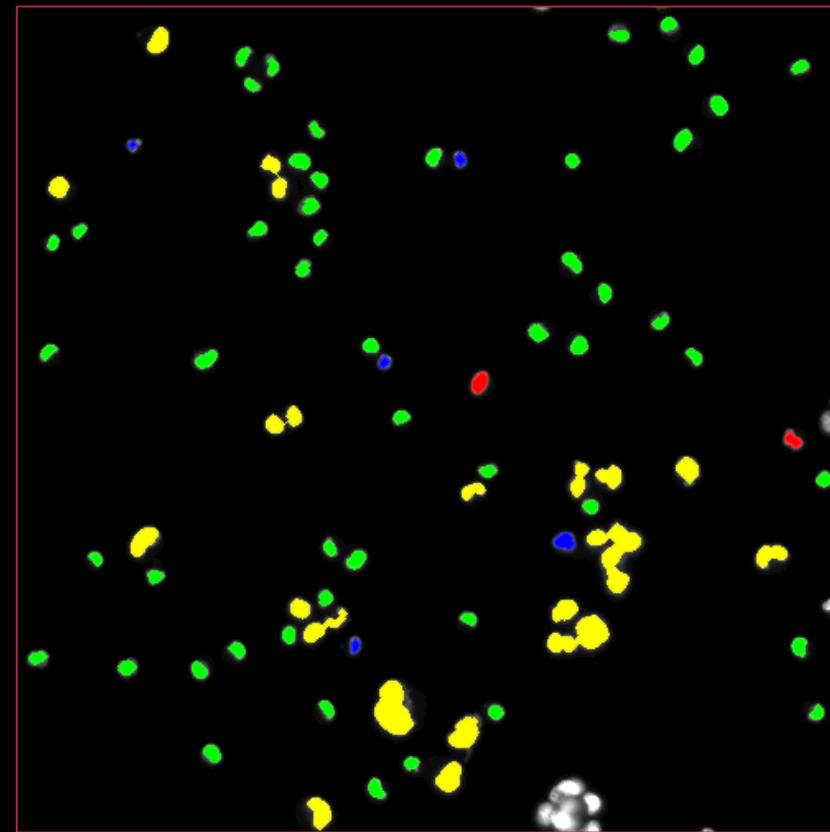
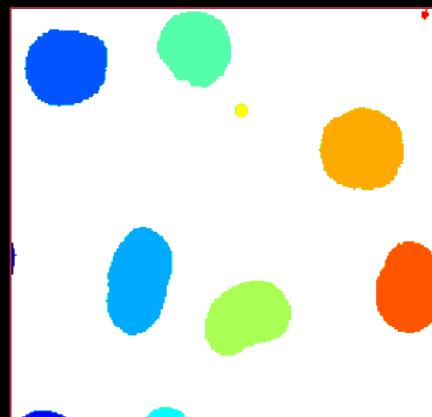
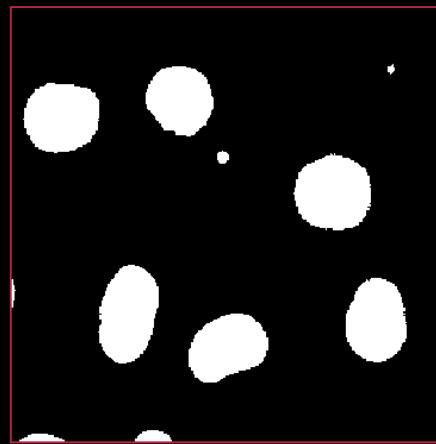
Tim B. Dyrby

DTU Compute

rapa@dtu.dk

<http://www.compute.dtu.dk/courses/02502>

Lecture 6 – BLOB analysis and feature based classification

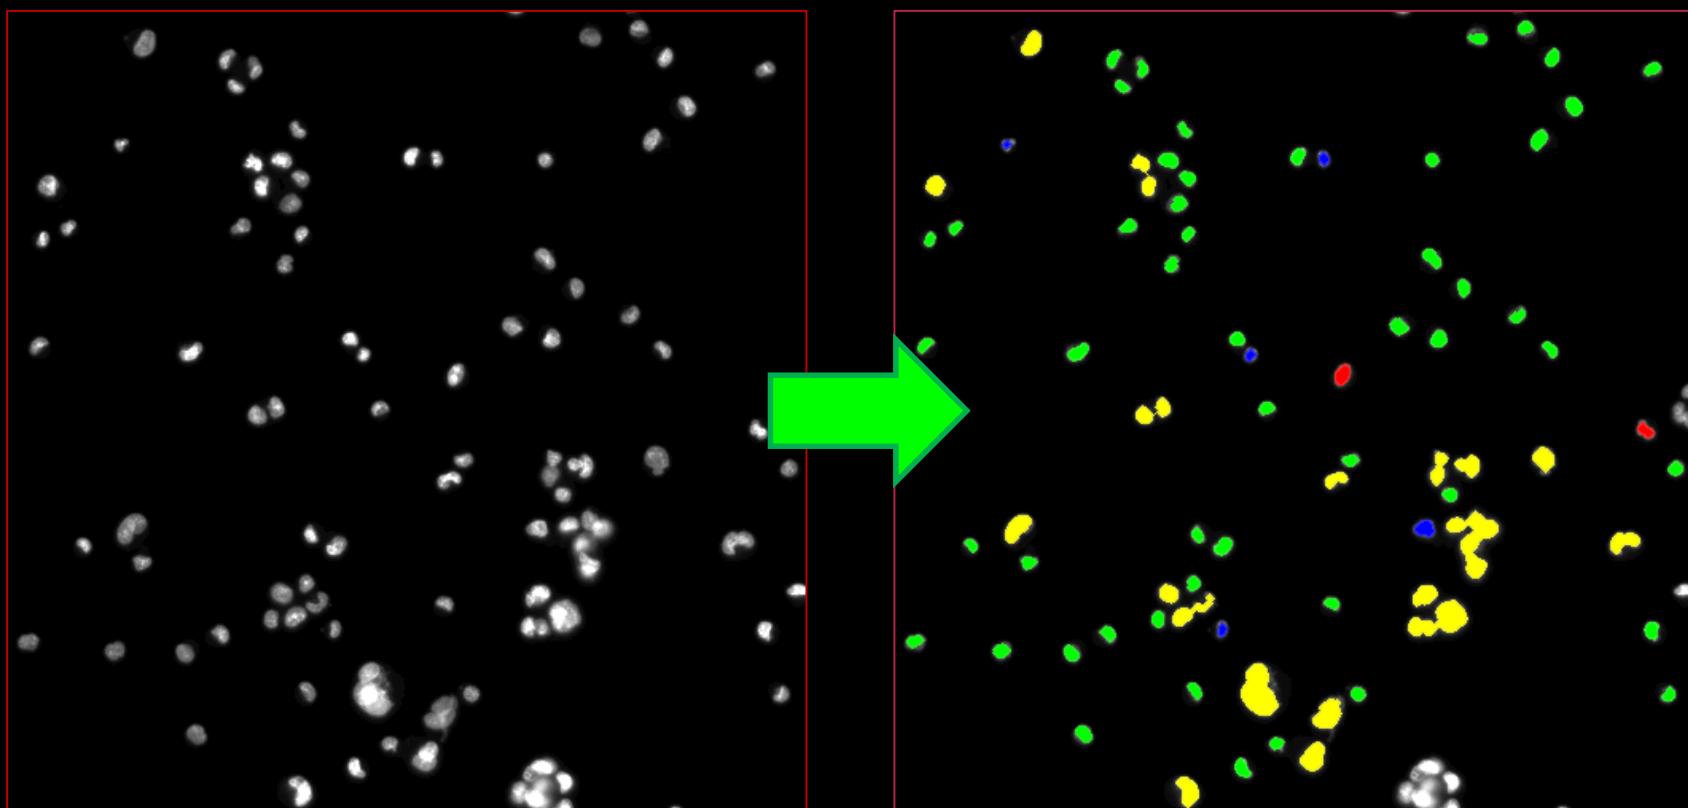


What can you do after today?

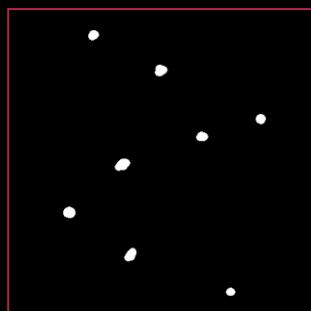
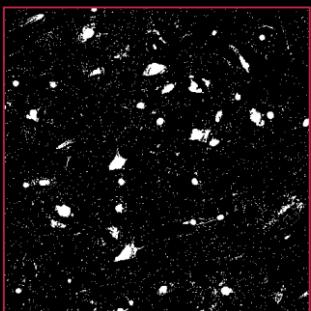
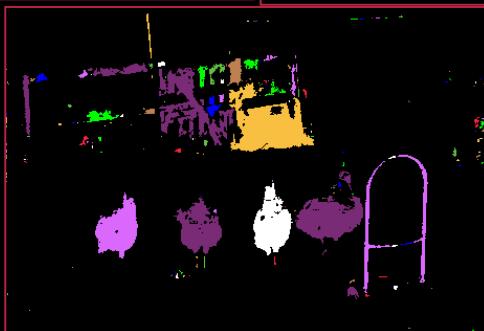
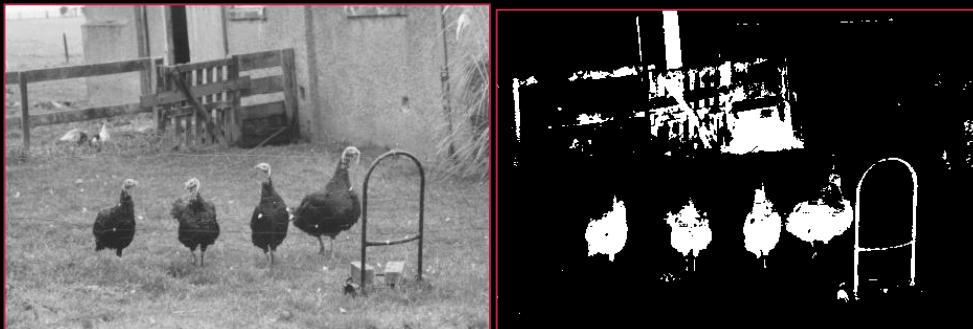
- Calculate the connected components of a binary image. Both using 4-connected and 8-connected neighbours
- Compute BLOB features including area, bounding box ratio, perimeter, center of mass, circularity, and compactness
- Describe a feature space
- Compute blob feature distances in feature space
- Classify binary objects based on their blob features
- Estimate feature value ranges using annotated training data
- Compute a confusion matrix
- Compute rates from a confusion matrix including sensitivity, specificity and accuracy
- Determine and discuss what is the importance of sensitivity and specificity given an image analysis problem

Object recognition

- Recognise objects in images
- Put them into different classes

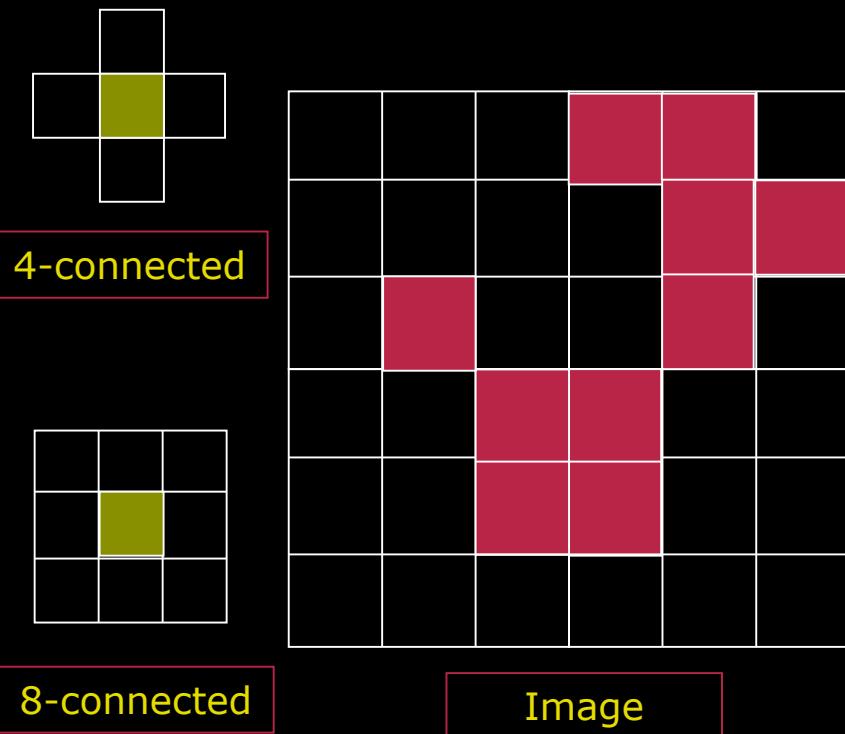


BLOB – what is it?



- BLOB = Binary Large Object
 - Group of connected pixels
- BLOB Analysis
 - *Connected component analysis*
 - *Object labelling*

Isolating a BLOB



■ What we want:

- For each object in the image, a list with its pixels

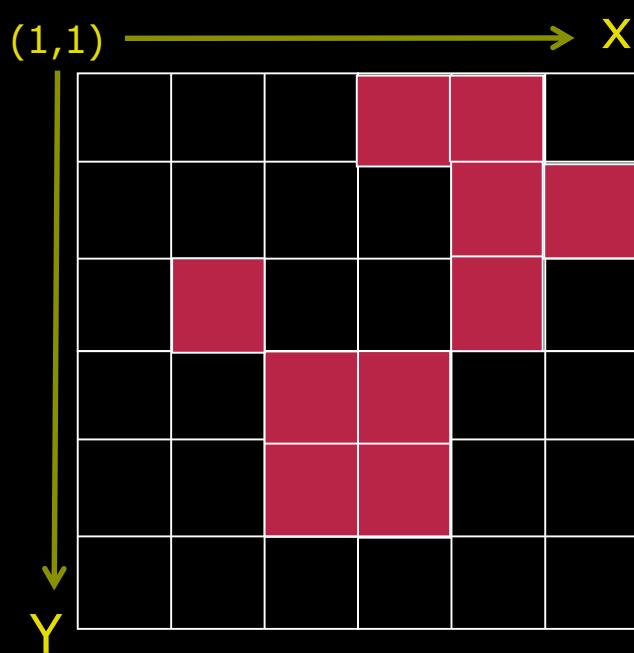
■ How do we get that?

- Connected component analysis

■ Connectivity

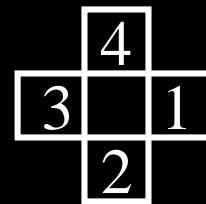
- Who are my neighbors?
- 4-connected
- 8-connected

Connected component analysis



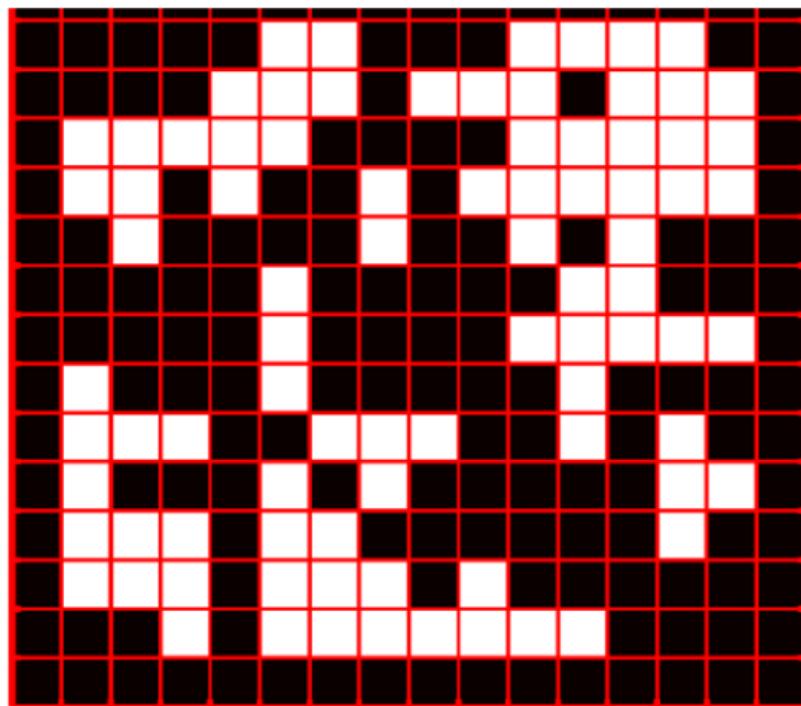
- Binary image
- Seed point: where do we start?
- *Grassfire* concept
 - Delete (burn) the pixels we visit
 - Visit all *connected* (4 or 8) neighbors

4-connected



BLOBs with 4- and 8- connectivity

A BLOB analysis is performed using both 4- and 8- connectivity. How many BLOBs are found using the two different connectivities?



3 and 7

9 and 5

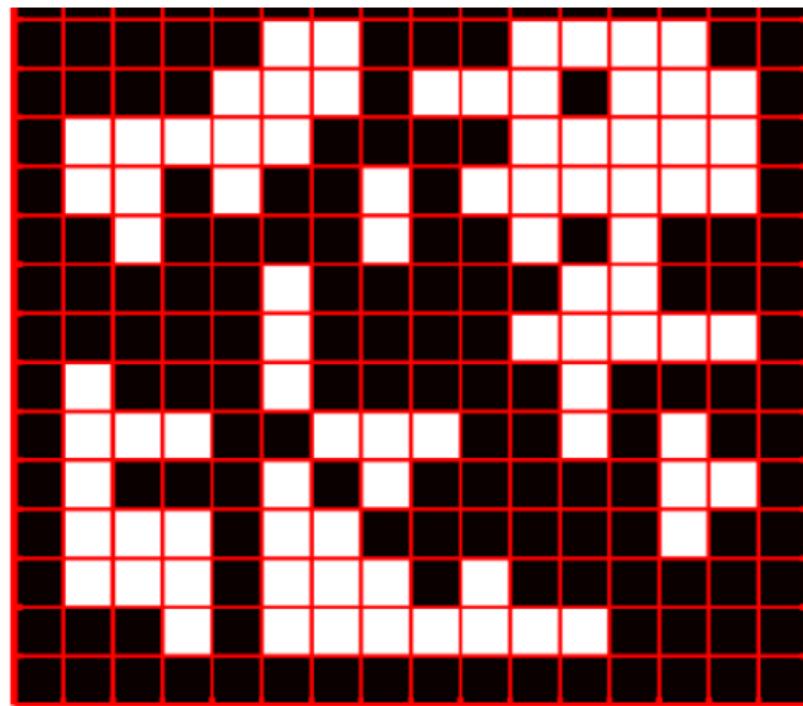
8 and 6

7 and 5

4 and 5

BLOBs with 4- and 8- connectivity

A BLOB analysis is performed using both 4- and 8- connectivity. How many BLOBS are found using the two different connectivities?



3 and 7

9 and 5

8 and 6

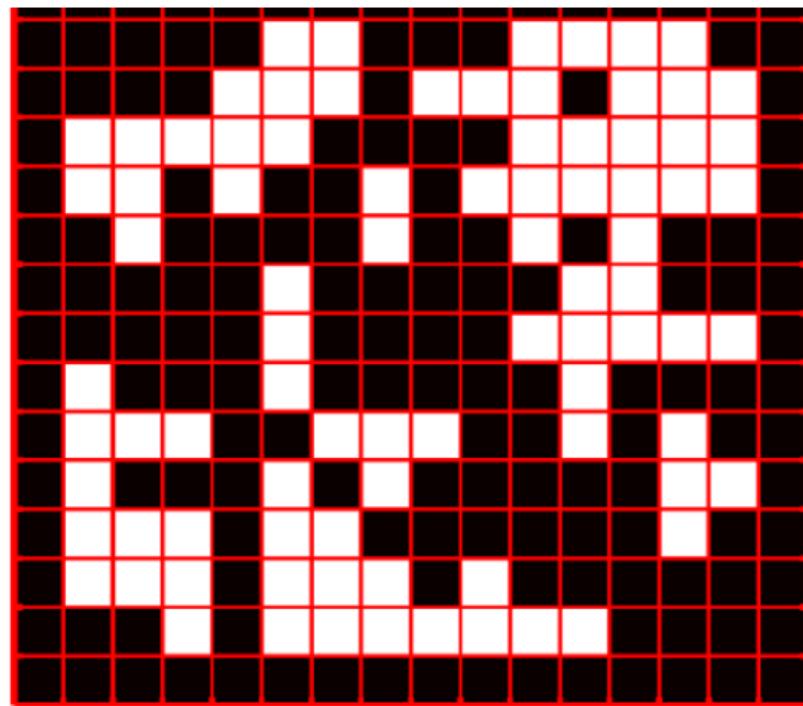
7 and 5

4 and 5

Start the presentation to see live content. For screen share software, share the entire screen. Get help at pollev.com/app

BLOBs with 4- and 8- connectivity

A BLOB analysis is performed using both 4- and 8- connectivity. How many BLOBS are found using the two different connectivities?



3 and 7

9 and 5

8 and 6

✓ 0%

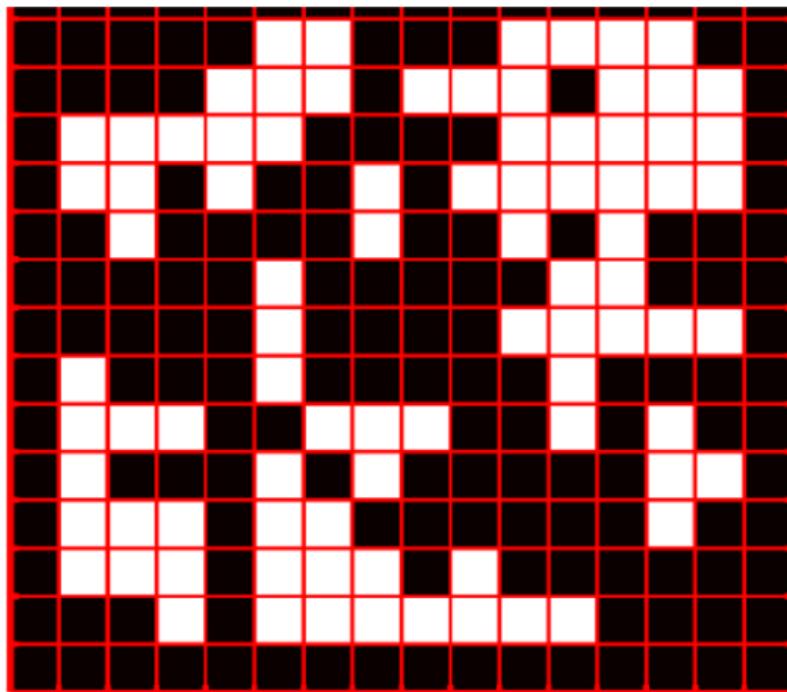
7 and 5

4 and 5

Start the presentation to see live content. For screen share software, share the entire screen. Get help at pollev.com/app

BLOBs with 4- and 8- connectivity

A BLOB analysis is performed using both 4- and 8- connectivity. How many BLOBs are found using the two different connectivities?



3 and 7

9 and 5

8 and 6

7 and 5

4 and 5

Start the presentation to see live content. For screen share software, share the entire screen. Get help at pollev.com/app

BLOB extraction and binary run length coding

A BLOB analysis using 4-connectivity is performed on the image seen in Figure 4. It is a 0-based (x,y) coordinate system with origin in the upper left corner. The largest BLOB is kept and the resulting image is coded using binary run-length coding. The code is:

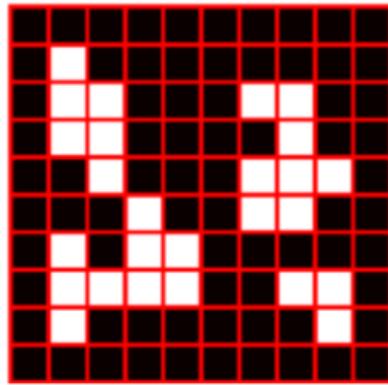


Figure 4: Binary image I. White pixels are foreground (1) and black pixels are background (0).

[5; (3; 3)]; [6; (1; 1)]; [6; (3; 4)]; [7; (1; 4)]; [8; (1; 1)]

[4; (3; 3)]; [6; (1; 2)]; [7; (3; 4)]; [7; (1; 4)]; [8; (1; 1)]

[5; (2; 3)]; [6; (1; 1)]; [6; (3; 4)]; [8; (1; 5)]; [8; (1; 1)]

[4; (3; 3)]; [6; (1; 2)]; [6; (3; 5)]; [7; (1; 4)]; [8; (1; 3)]

[5; (2; 3)]; [6; (1; 1)]; [6; (3; 4)]; [7; (2; 4)]; [8; (2; 3)]

BLOB extraction and binary run length coding

A BLOB analysis using 4-connectivity is performed on the image seen in Figure 4. It is a 0-based (x,y) coordinate system with origin in the upper left corner. The largest BLOB is kept and the resulting image is coded using binary run-length coding. The code is:

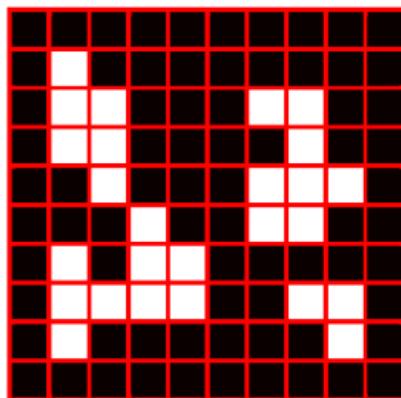


Figure 4: Binary image I. White pixels are foreground (1) and black pixels are background (0).

[5; (3; 3)]; [6; (1;
1)]; [6; (3; 4)]; [7;
(1; 4)]; [8; (1; 1)]

[4; (3; 3)]; [6; (1;
2)]; [7; (3; 4)]; [7;
(1; 4)]; [8; (1; 1)]

[5; (2; 3)]; [6; (1;
1)]; [6; (3; 4)]; [8;
(1; 5)]; [8; (1; 1)]

[4; (3; 3)]; [6; (1;
2)]; [6; (3; 5)]; [7;
(1; 4)]; [8; (1; 3)]

[5; (2; 3)]; [6; (1;
1)]; [6; (3; 4)]; [7;
(2; 4)]; [8; (2; 3)]

BLOB extraction and binary run length coding

A BLOB analysis using 4-connectivity is performed on the image seen in Figure 4. It is a 0-based (x,y) coordinate system with origin in the upper left corner. The largest BLOB is kept and the resulting image is coded using binary run-length coding. The code is:

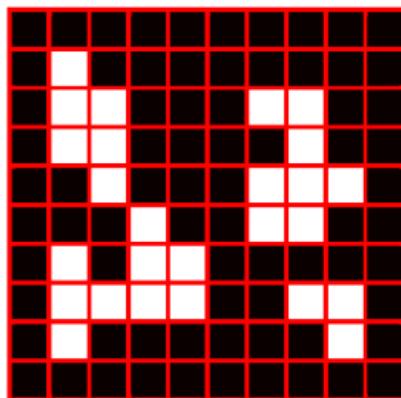


Figure 4: Binary image I. White pixels are foreground (1) and black pixels are background (0).

[5; (3; 3)]; [6; (1; 1)]; [6; (3; 4)]; [7; (1; 4)]; [8; (1; 1)]

✓ 0%

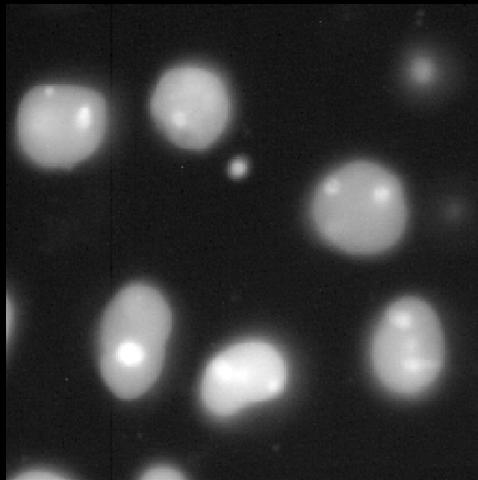
[4; (3; 3)]; [6; (1; 2)]; [7; (3; 4)]; [7; (1; 4)]; [8; (1; 1)]

[5; (2; 3)]; [6; (1; 1)]; [6; (3; 4)]; [8; (1; 5)]; [8; (1; 1)]

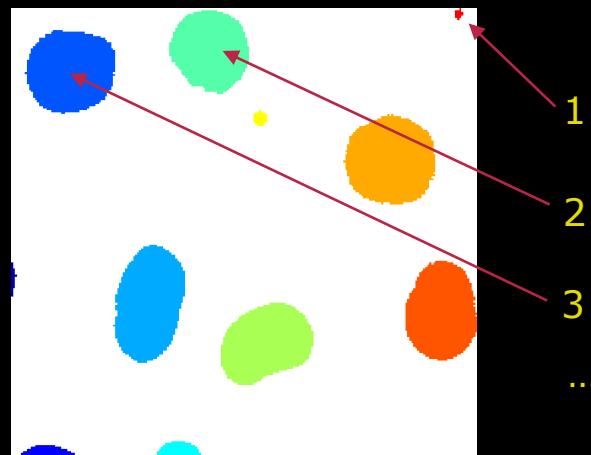
[4; (3; 3)]; [6; (1; 2)]; [6; (3; 5)]; [7; (1; 4)]; [8; (1; 3)]

[5; (2; 3)]; [6; (1; 1)]; [6; (3; 4)]; [7; (2; 4)]; [8; (2; 3)]

The result of connected component analysis



- An image where each BLOB (component) is labelled
- Each blob now has a unique ID number
- What do we do with these blobs?



Features



- Feature
 - A prominent or distinctive aspect, quality, or characteristic
 - *This radio has many good features*
- Car (Ford-T) features
 - 4 wheels
 - 2 doors
 - 540 kg
 - 20 hp

Feature vector



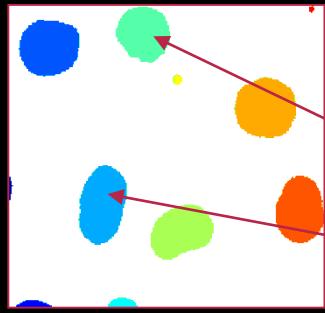
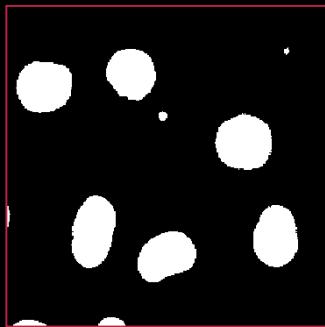
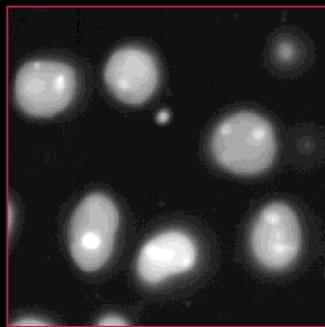
$f=[4, 2, 540, 20]$

- Feature vector
 - Vector with all the features for one object
- Ford-T features
 - 4 wheels
 - 2 doors
 - 540 kg
 - 20 hp
- Ford Fiesta features
 - 4 wheels
 - 3 doors
 - 1100 kg
 - 90 hp



$f=[4, 3, 1100, 90]$

Feature extractions



- Compute features for each BLOB that can be used to identify it
 - Size
 - Shape
 - Position
- From image operations to mathematical operations
 - **Input:** a list of pixel positions
 - **Output:** Feature vector
- First step: remove invalid BLOBS
 - too small or big- using morphological operations for example
 - border BLOBS

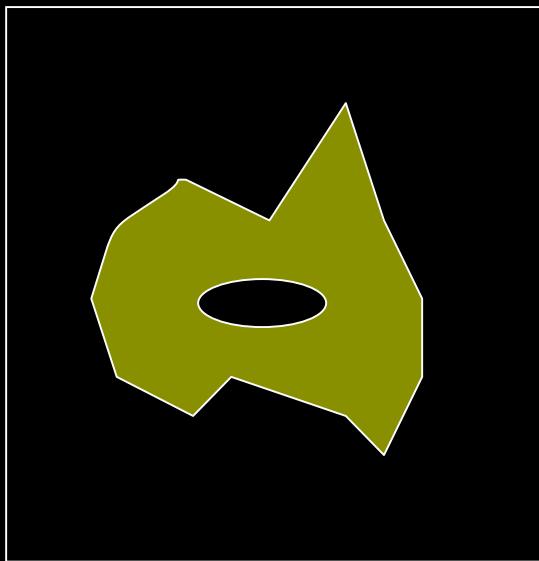
Feature vector = [2,1,...,3]

Feature vector = [4,7,...,0]

BLOB Features

■ Area

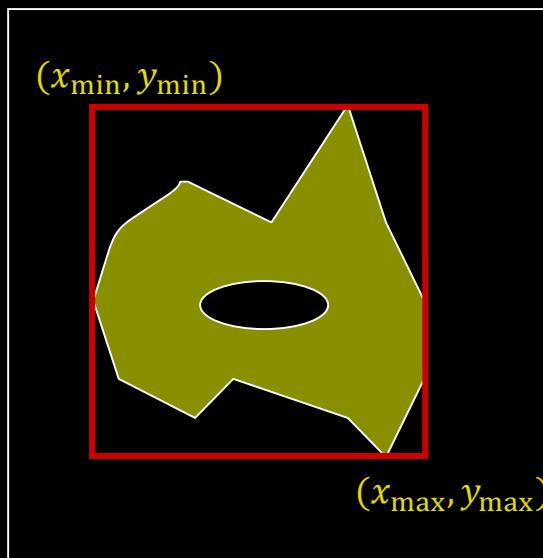
- number of pixels in the BLOB
- Can be used to remove noise (small BLOBS)



One BLOB

BLOB Features

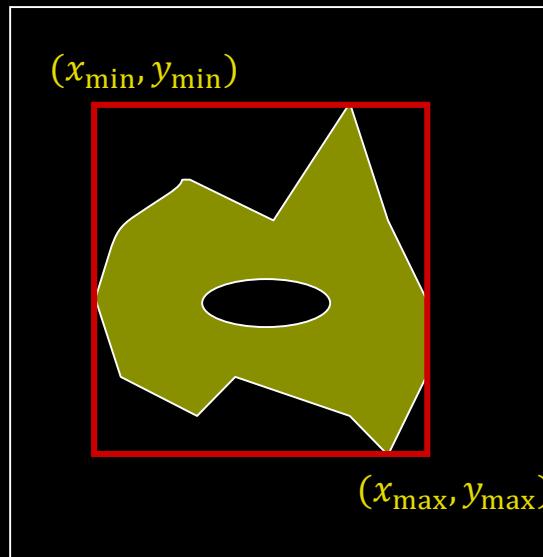
- Bounding box
 - Minimum rectangle that contains the BLOB
 - Height: $y_{\max} - y_{\min}$
 - Width: $x_{\max} - x_{\min}$
 - Bounding box ratio:
$$\frac{y_{\max} - y_{\min}}{x_{\max} - x_{\min}}$$
 - tells if the BLOB is elongated



One BLOB

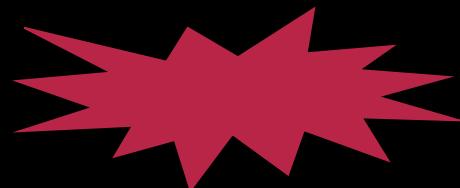
BLOB Features

- Bounding box
 - Bounding box area:



One BLOB

$$\text{Compactness} = \frac{\text{BLOB Area}}{(y_{\max} - y_{\min}) \cdot (x_{\max} - x_{\min})}$$



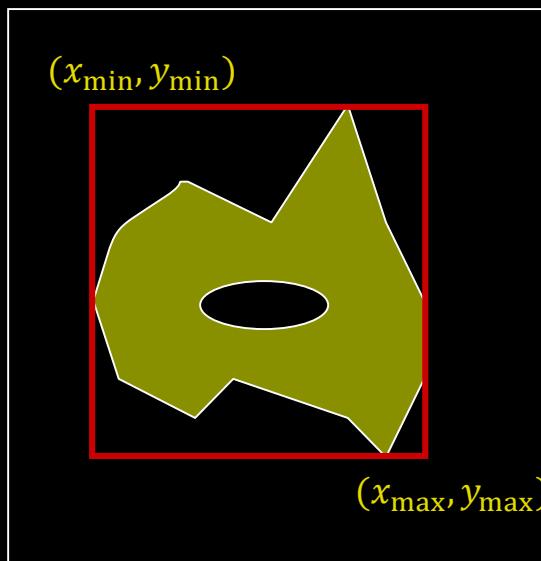
Not compact



Compact

BLOB Features

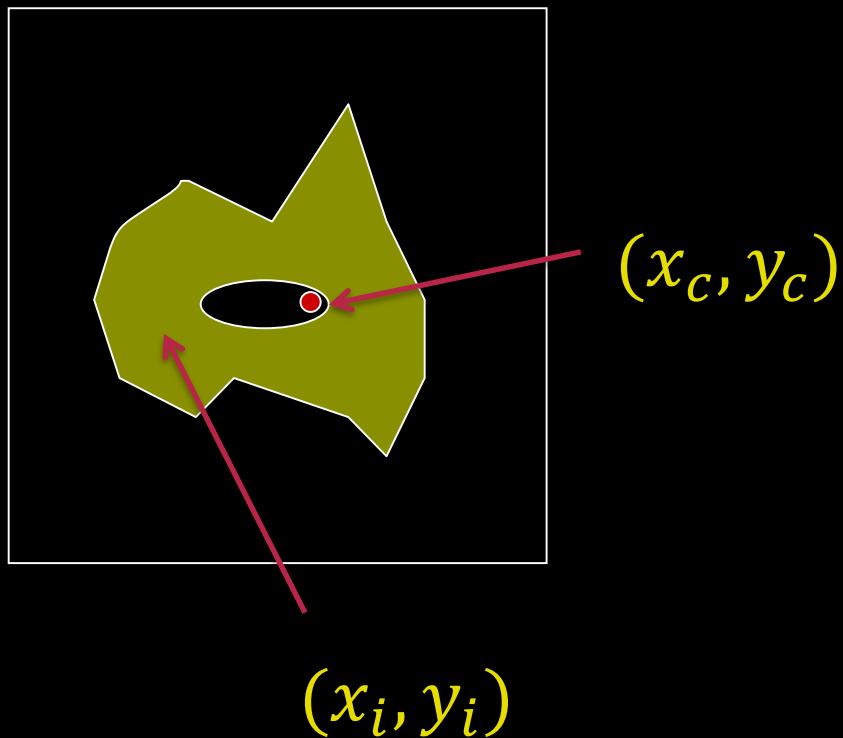
- Bounding box ratio
 - Bounding box height divided by the width



One BLOB

BLOB Features

- Center of mass (x_c, y_c)

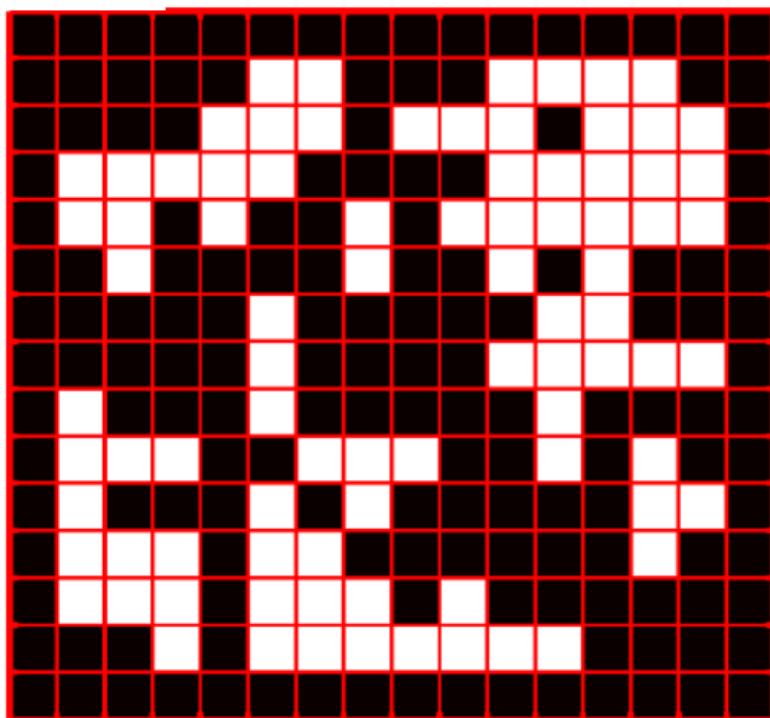


$$x_c = \frac{1}{N} \sum_{i=1}^N x_i$$

$$y_c = \frac{1}{N} \sum_{i=1}^N y_i$$

BLOB Center of Mass

The smallest BLOB is found using 4-connectivity. What is the center of mass of this BLOB. The image has origin (0,0) and uses a (x,y) coordinate system.



(12, 1.5)

(5, 8.5)

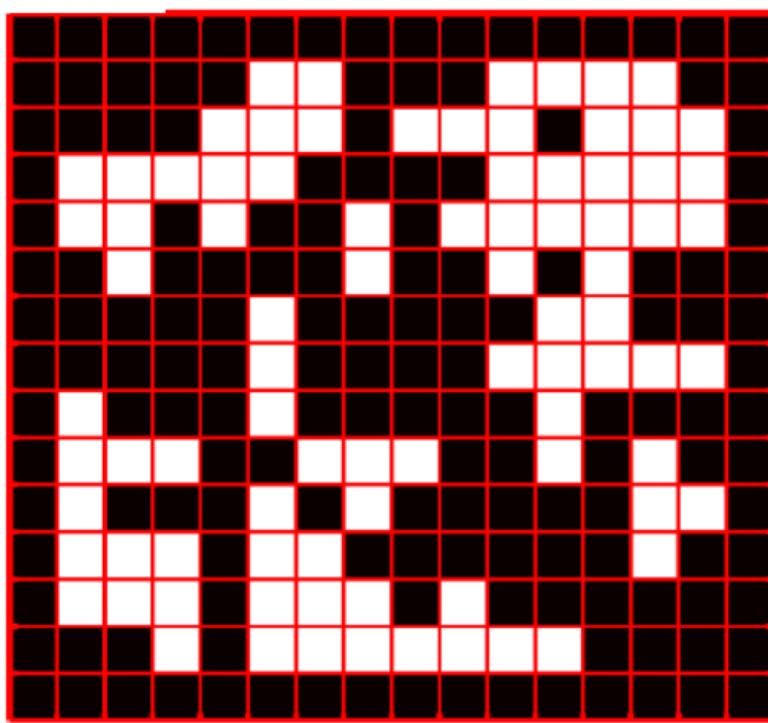
(6.5, 3.5)

(4.5, 0.5)

(7, 4.5)

BLOB Center of Mass

The smallest BLOB is found using 4-connectivity. What is the center of mass of this BLOB. The image has origin (0,0) and uses a (x,y) coordinate system.



(12, 1.5)

(5, 8.5)

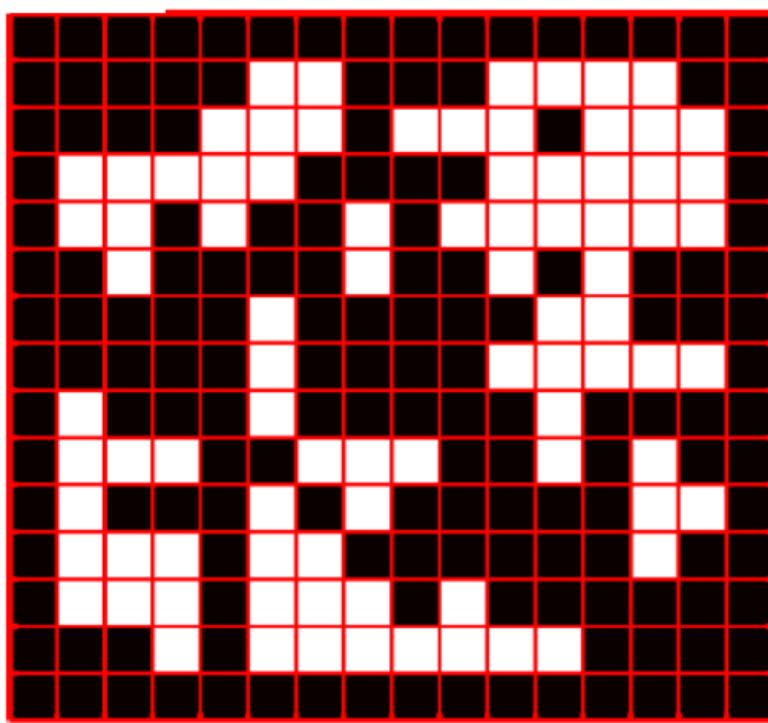
(6.5, 3.5)

(4.5, 0.5)

(7, 4.5)

BLOB Center of Mass

The smallest BLOB is found using 4-connectivity. What is the center of mass of this BLOB. The image has origin (0,0) and uses a (x,y) coordinate system.



(12, 1.5)

(5, 8.5)

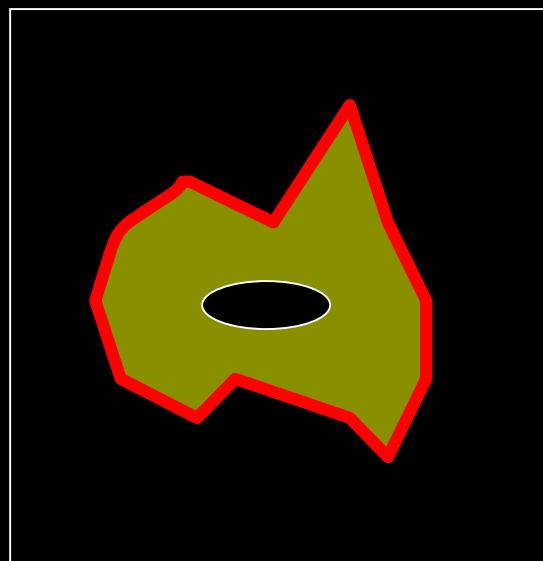
(6.5, 3.5)

(4.5, 0.5)

(7, 4.5)

✓ 0%

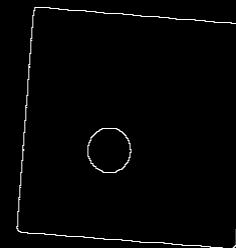
BLOB Features



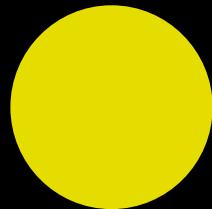
One BLOB

- Perimeter
 - Length of perimeter
 - How can we compute that?
- In practice (in Matlab) it is computed differently and more accurately

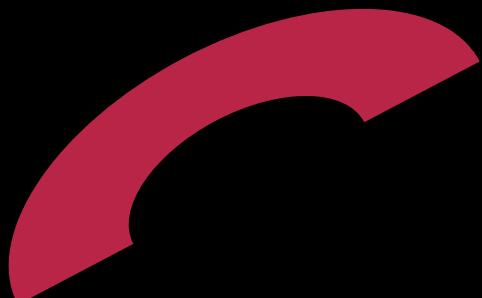
$$\sum ((f(x, y) \oplus SE) - f(x, y))$$



BLOB Features - circularity



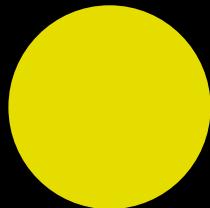
Circle like



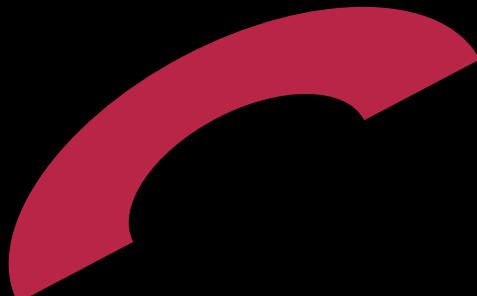
Not circle like

- How much does it look like a circle?
- Circle
 - Area $A = \pi r^2$
 - Perimeter $P = 2\pi r$
- New object assumed to be a circle
 - Measured perimeter P_m
 - Measured area A_m
- Estimate perimeter from (measured) area
 - Estimated perimeter $P_e = 2\sqrt{\pi A_m}$

BLOB Features - circularity



Circle like



Not circle like

- Compare the perimeters
 - Measured perimeter P_m
 - Estimated perimeter $P_e = 2\sqrt{\pi A_m}$
- Circularity 1:

$$\text{Circularity} = \frac{P_m}{P_e} = \frac{P_m}{2\sqrt{\pi A_m}}$$

Circularity math



$$\text{Circularity} = \frac{P_m}{P_e} = \frac{P_m}{2\sqrt{\pi A_m}}$$

$P_m < P_e$

$P_m = P_e$

$P_m > P_e$

Circularity math



$$\text{Circularity} = \frac{P_m}{P_e} = \frac{P_m}{2\sqrt{\pi A_m}}$$

$$P_m < P_e$$

$$P_m = P_e$$

$$P_m > P_e$$

Start the presentation to see live content. For screen share software, share the entire screen. Get help at pollev.com/app

Circularity math



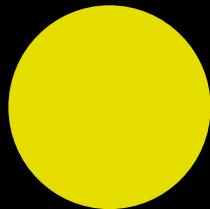
$$\text{Circularity} = \frac{P_m}{P_e} = \frac{P_m}{2\sqrt{\pi A_m}}$$

$$P_m = P_e$$

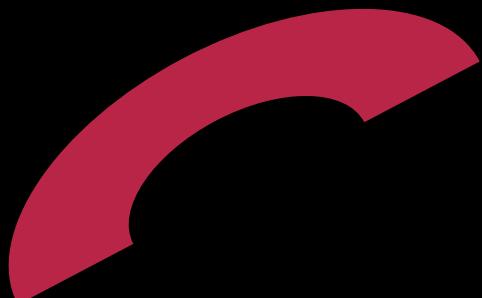
$$P_m > P_e$$

✓ 0%

BLOB Features - circularity



Circle like



Not circle like

- Compare the perimeters

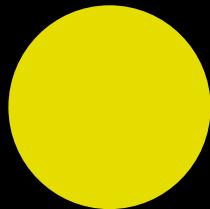
- Measured perimeter P_m
- Estimated perimeter $P_e = 2\sqrt{\pi A_m}$

- Circularity:

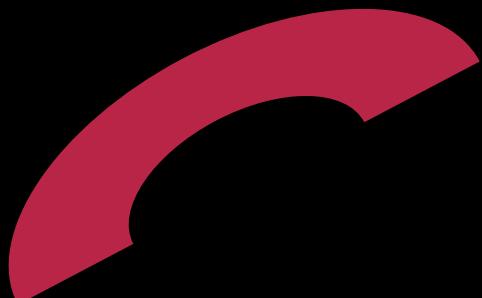
$$\text{Circularity} = \frac{P_m}{P_e} = \frac{P_m}{2\sqrt{\pi A_m}}$$

- This measure will normally be ≥ 1

BLOB Features – circularity inverse



Circle like



Not circle like

- Compare the perimeters

- Measured perimeter P_m
- Estimated perimeter $P_e = 2\sqrt{\pi A_m}$

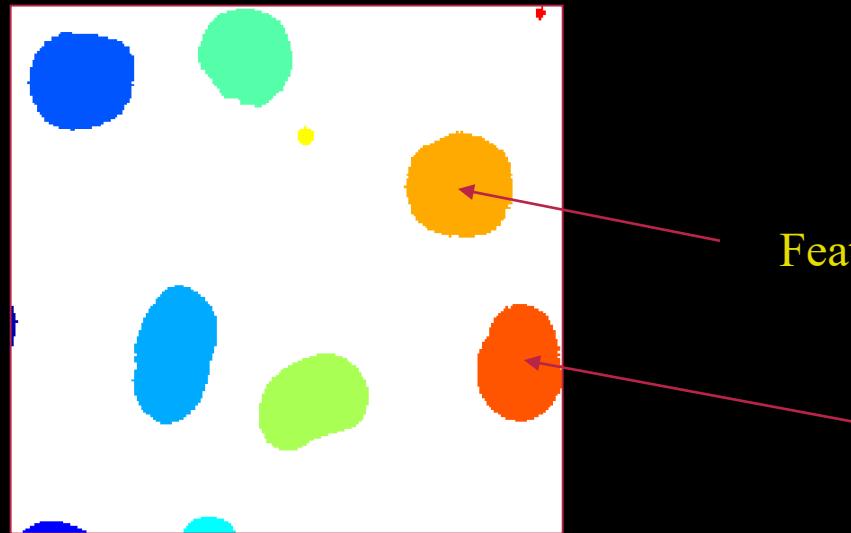
- Circularity (inverse):

$$\text{Circularity inverse} = \frac{P_e}{P_m} = \frac{2\sqrt{\pi A_m}}{P_m}$$

- This measure will normally be ≤ 1

After feature extraction

Area, compactness, circularity etc calculated for all BLOB



Feature vector = [2,1,...,3]

Feature vector = [4,7,...,0]

One feature vector per blob

BLOB Classification

■ Classification

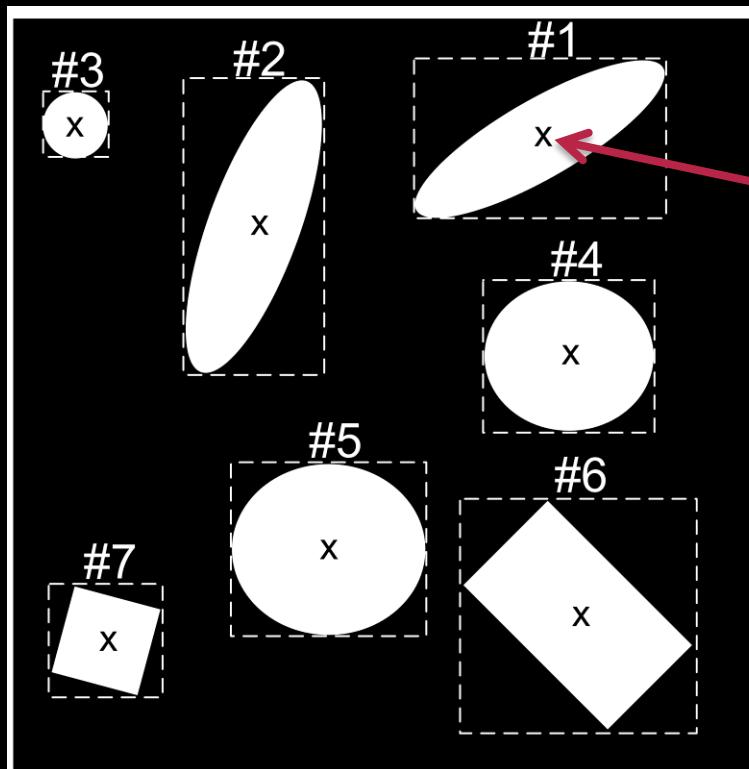
- Put a BLOB into a *class*

■ *Classes* are normally pre-defined

- *Car*
- *Bus*
- *Motorcycle*
- *Scooter*

■ Object recognition

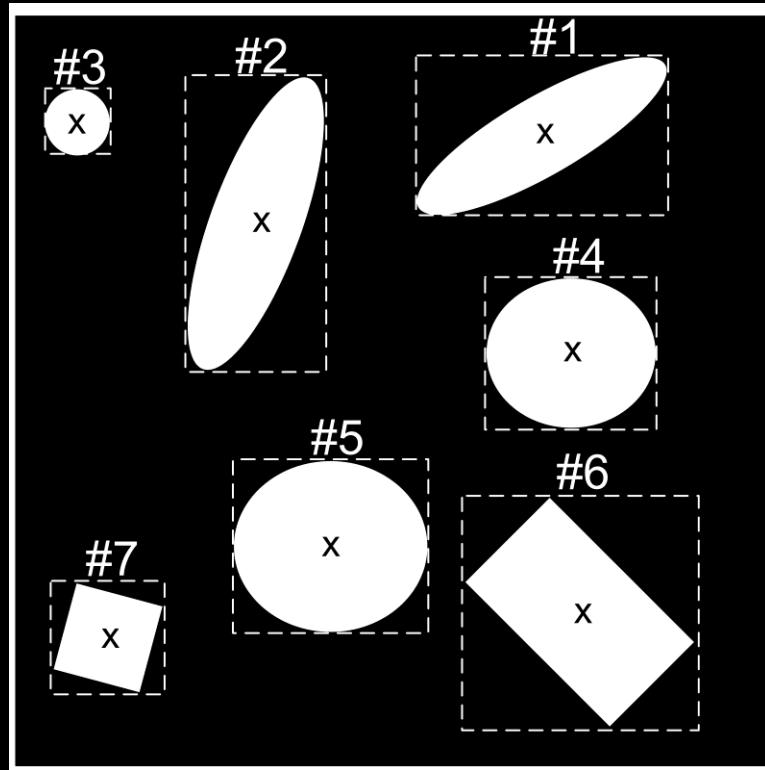
Object recognition: Circle example



BLOB number	Circu-larity	Area (pixels)
1	0.31	6561
2	0.40	6544
3	0.98	890
4	0.97	6607
5	0.99	6730
6	0.52	6611
7	0.75	2073

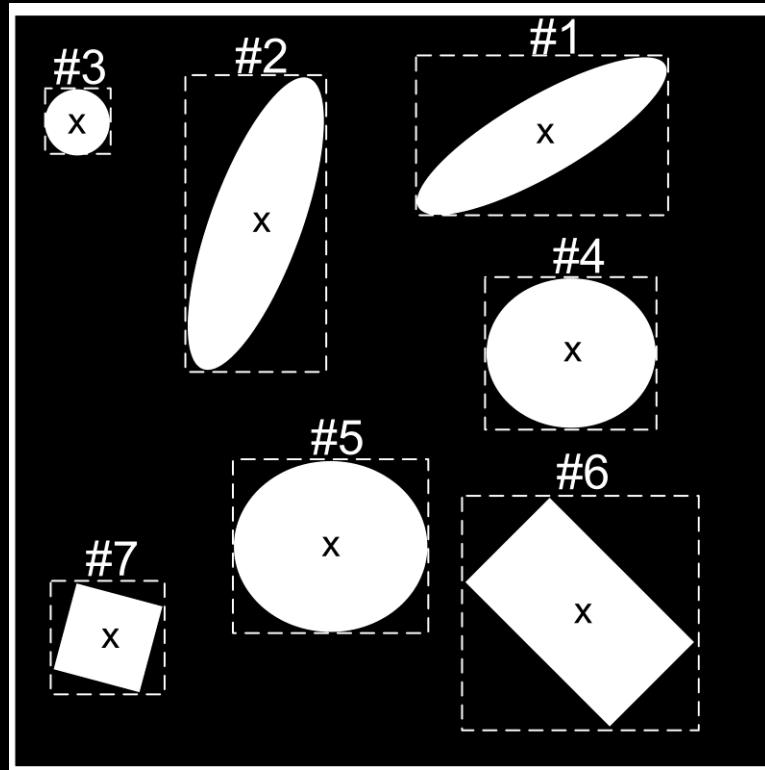
Which objects are circles?

Circle classification



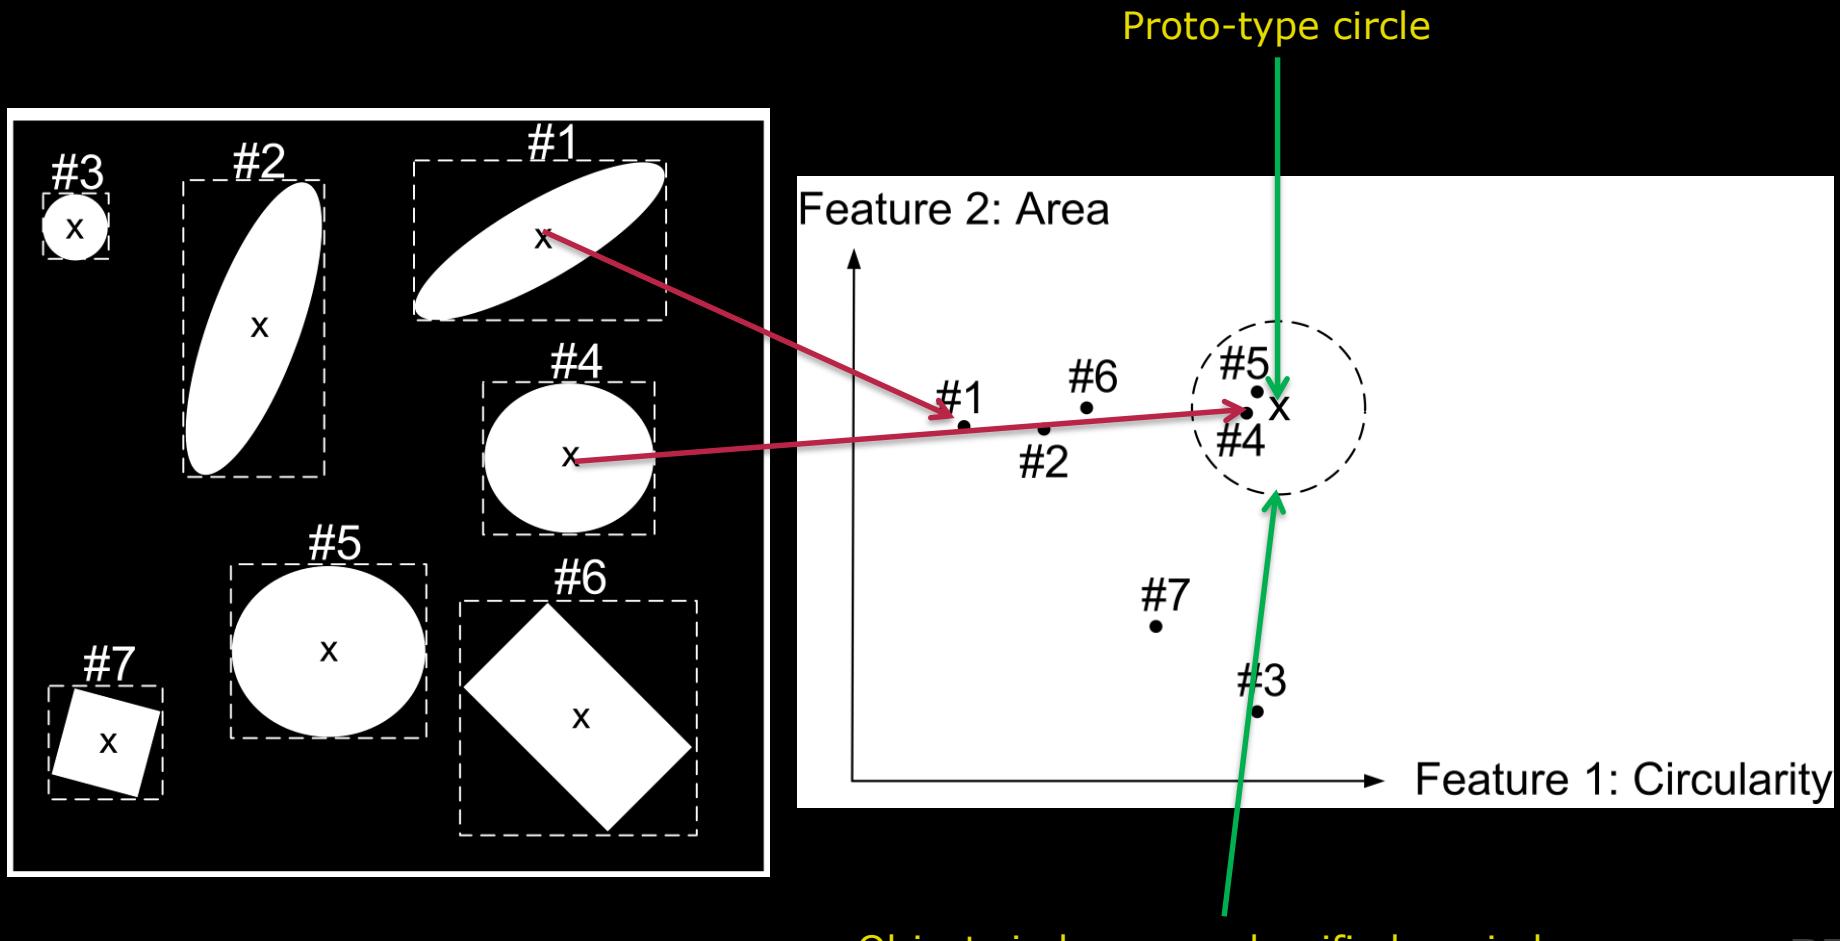
- Two classes:
 - Circle
 - Not-circle
- Lets make a model of a *proto-type* circle

Circle classification

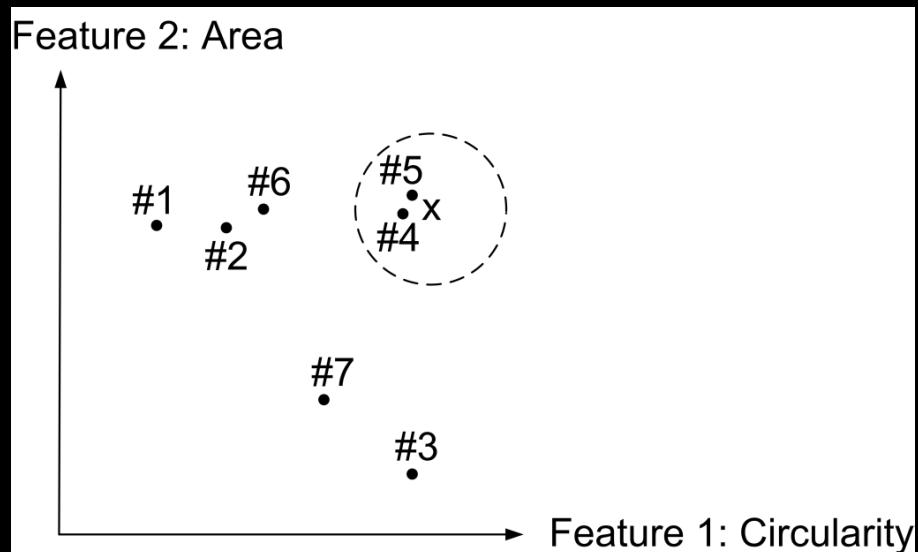


- Proto-type circle
 - Circularity : 1
 - Area: 6700

Feature Space

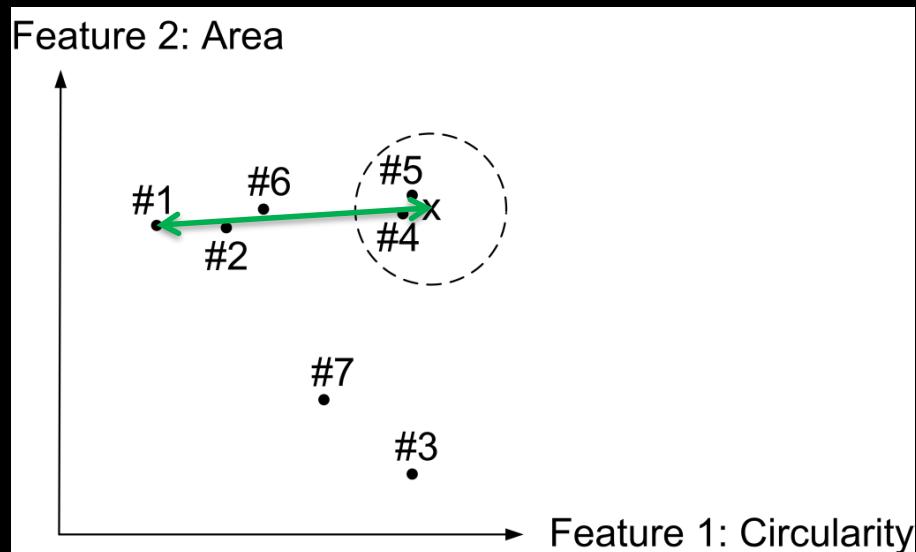


Feature space



- Proto-type circle
 - Circularity : 1
 - Area: 6700
- Some slack is added to allow non-perfect circles
 - Circularity: 1 +/- 0.15

Feature space - distances



- How do we decide if an object is inside the circle?
- Feature space distance
- Euclidean distance in features space

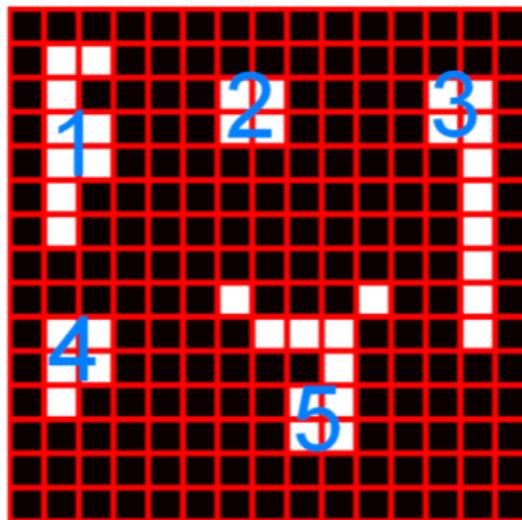
Blob 1: circularity: 0.31, Area : 6561

$$D = \sqrt{(0.31 - 1)^2 + (6561 - 6700)^2}$$

Dominate all! – normalisation needed

BLOB Classification

A BLOB analysis using 8-connectivity has been performed on the image seen in Figure 12 and the five found BLOBs have been marked with numbers. The BLOB features *area* and *compactness* have been computed for the five BLOBs. A reference BLOB has an area of 10 pixels and a compactness of 0.5. The Euclidean distance in feature space has been computed between the five BLOBs and the reference BLOB. Which of the five BLOBs has the minimum distance?



1

2

3

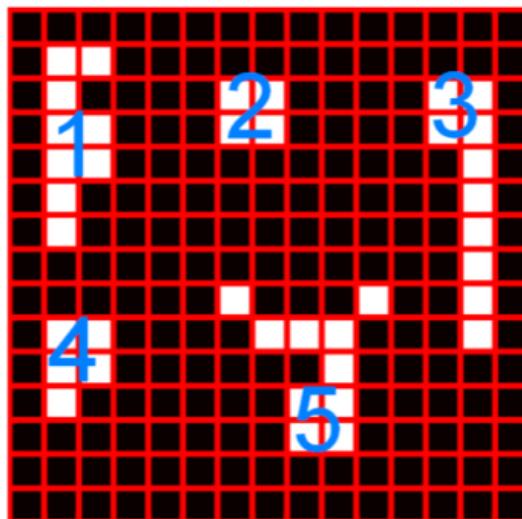
4

5

Start the presentation to see live content. For screen share software, share the entire screen. Get help at pollev.com/app

BLOB Classification

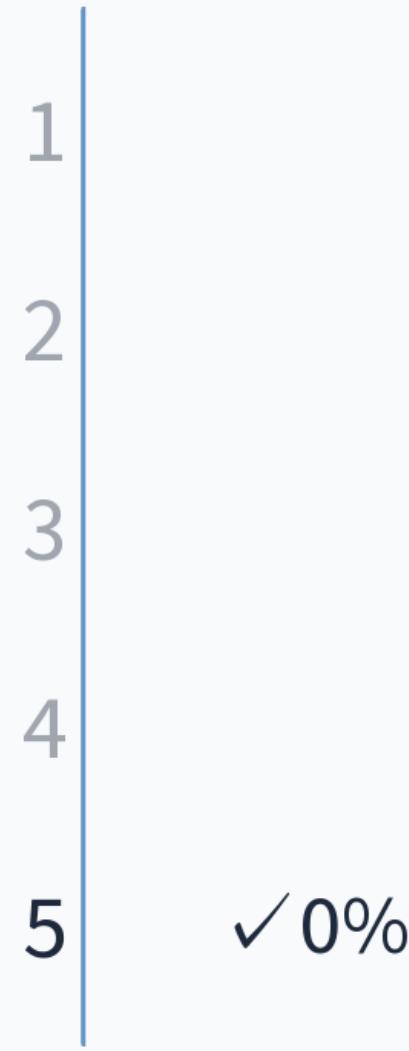
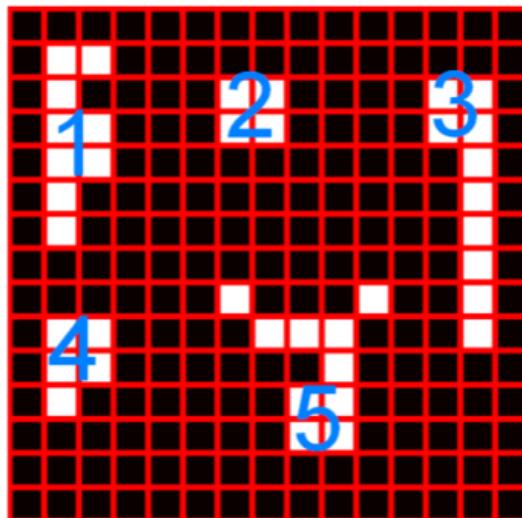
A BLOB analysis using 8-connectivity has been performed on the image seen in Figure 12 and the five found BLOBS have been marked with numbers. The BLOB features *area* and *compactness* have been computed for the five BLOBS. A reference BLOB has an area of 10 pixels and a compactness of 0.5. The Euclidean distance in feature space has been computed between the five BLOBS and the reference BLOB. Which of the five BLOBS has the minimum distance?



Start the presentation to see live content. For screen share software, share the entire screen. Get help at pollev.com/app

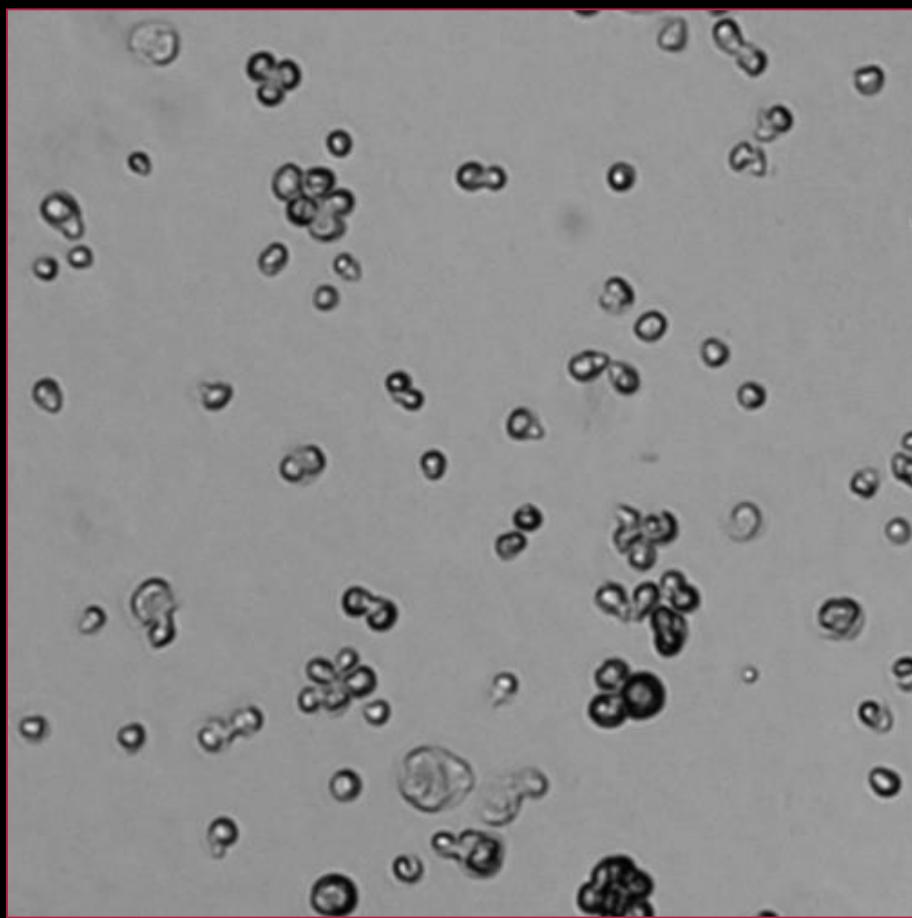
BLOB Classification

A BLOB analysis using 8-connectivity has been performed on the image seen in Figure 12 and the five found BLOBS have been marked with numbers. The BLOB features *area* and *compactness* have been computed for the five BLOBS. A reference BLOB has an area of 10 pixels and a compactness of 0.5. The Euclidean distance in feature space has been computed between the five BLOBS and the reference BLOB. Which of the five BLOBS has the minimum distance?

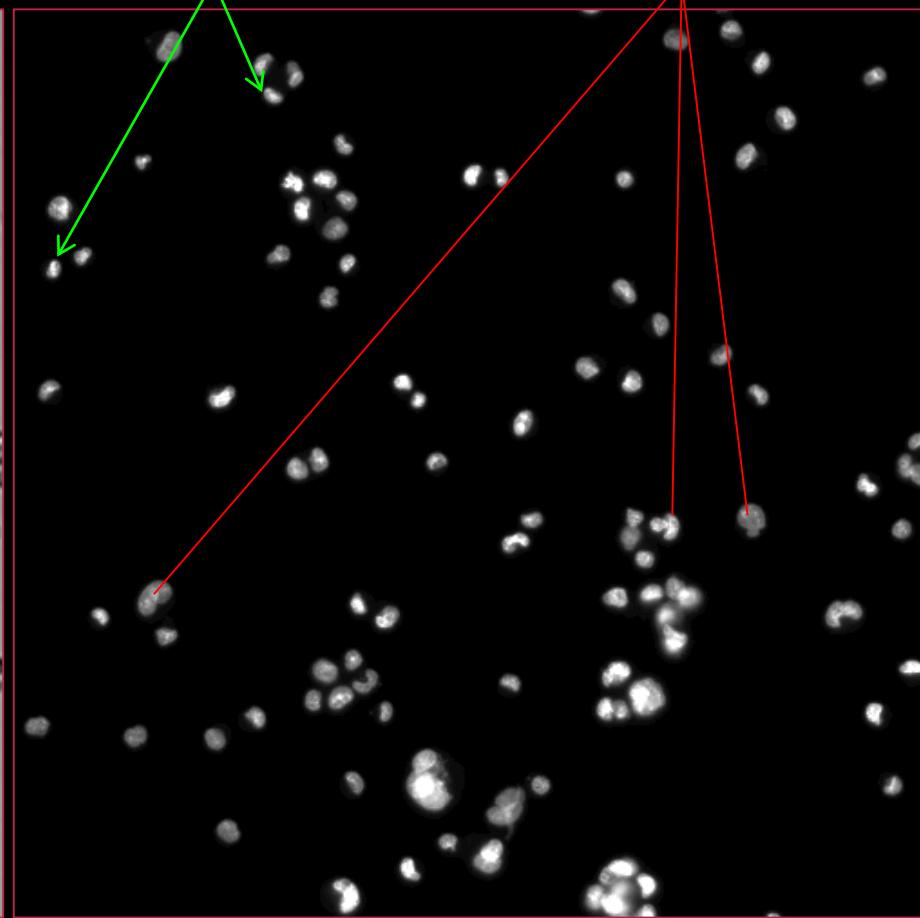


Start the presentation to see live content. For screen share software, share the entire screen. Get help at pollev.com/app

Cell classification



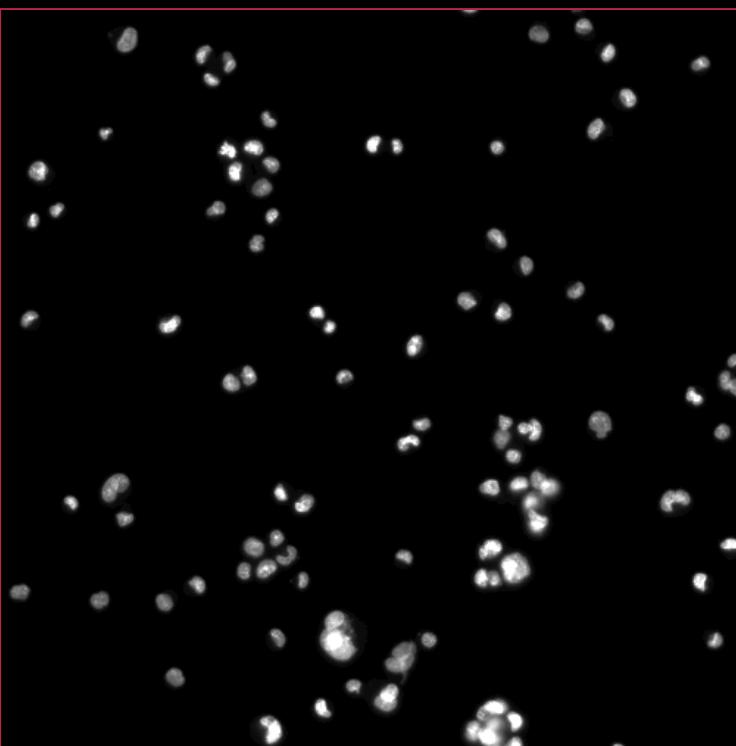
UV Microscopy



Fluorescence Microscopy (DAPI)

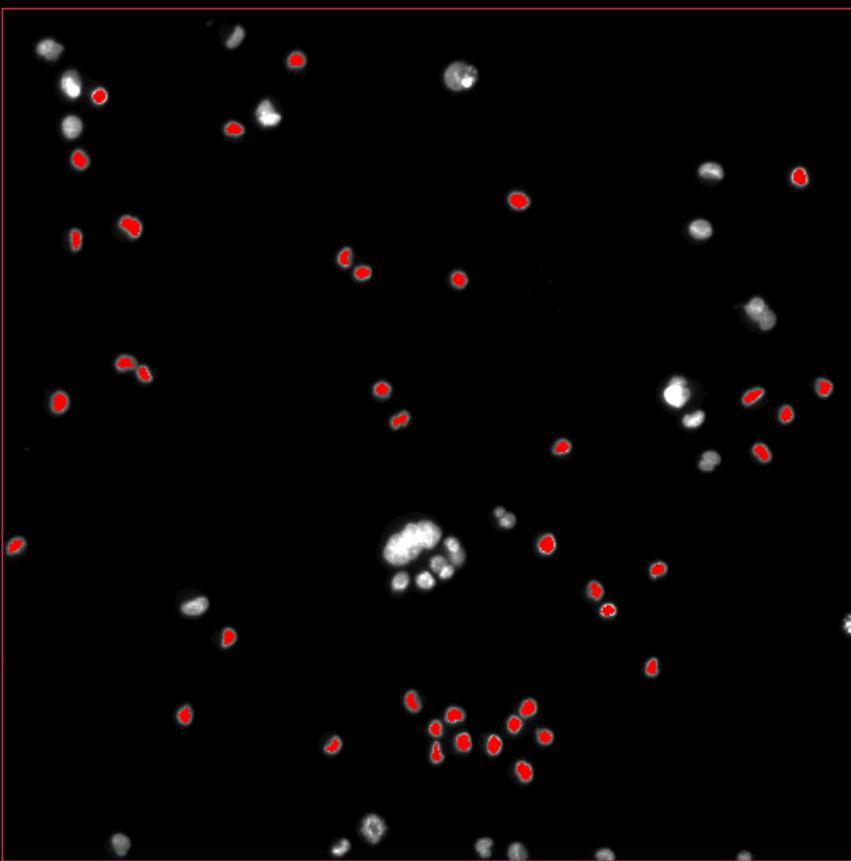
Images from ChemoMetec A/S

Nuclei classification



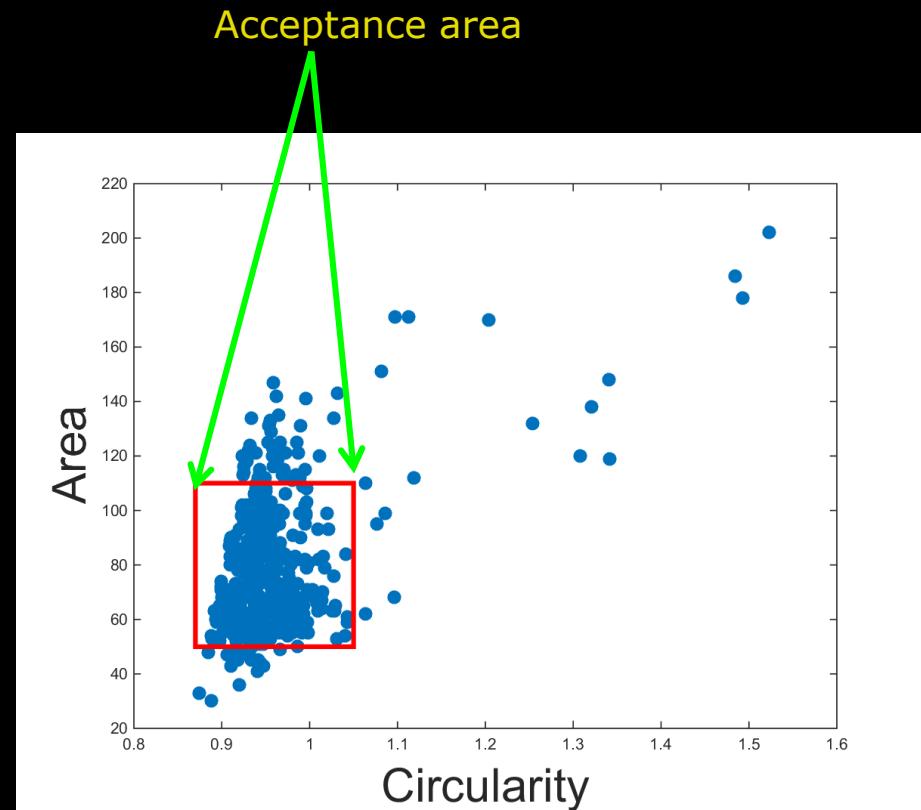
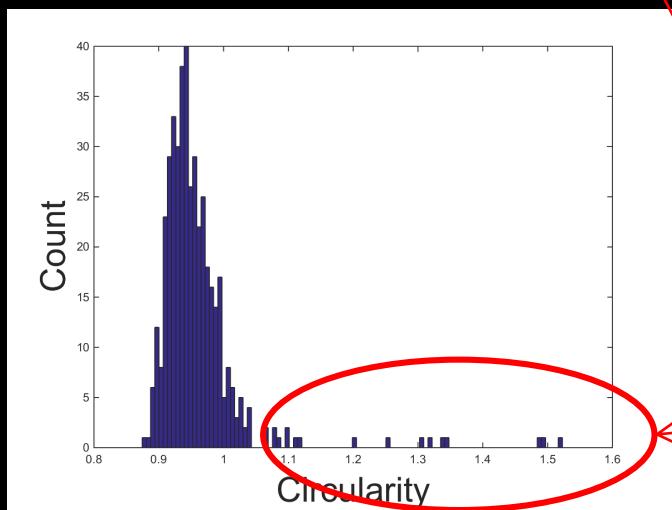
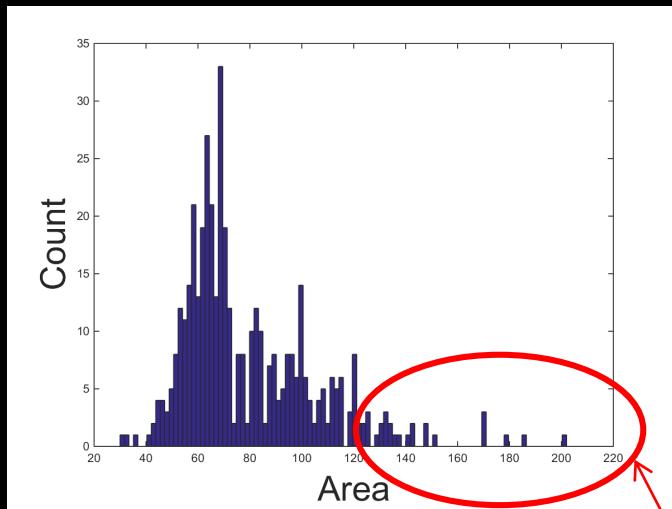
- DAPI image
- Two classes
 - Single nuclei
 - Noise
 - Multiple nuclei together
 - Debris
 - Other noise

Training and annotation



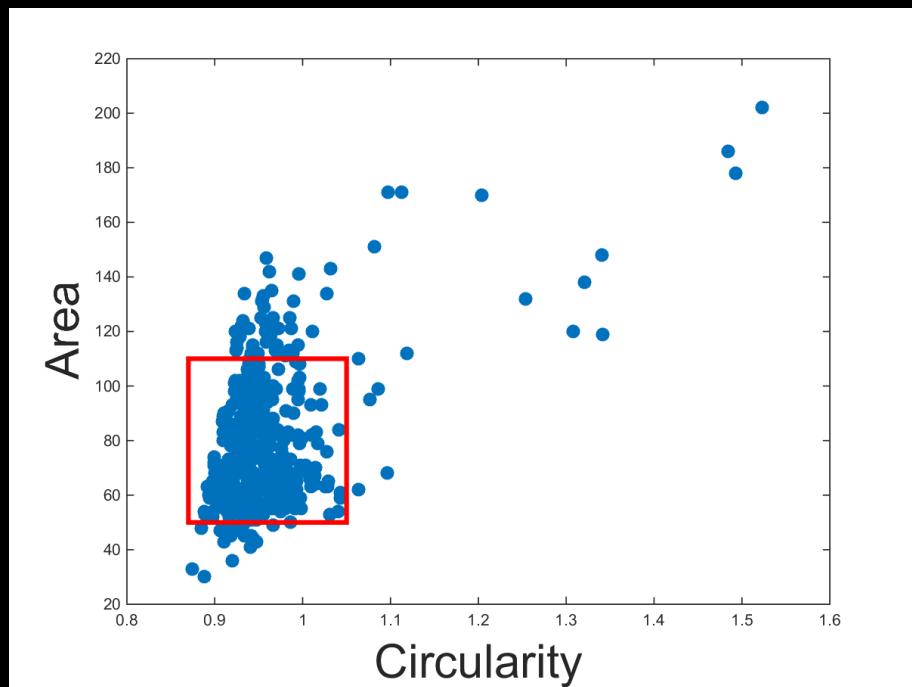
- Selection of true single nuclei marked
- Thresholding
- BLOB Analysis
 - Circularity
 - Area

Training data - analysis



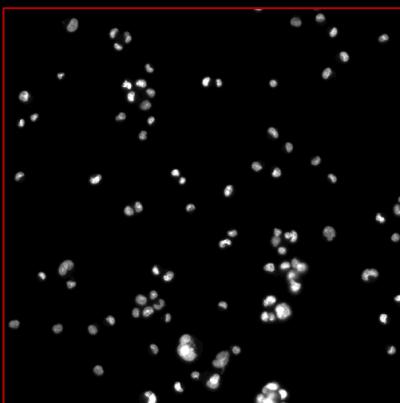
Probably outliers

Feature ranges



Feature	Min	Max
Area	50	110
Circularity	0.87	1.05

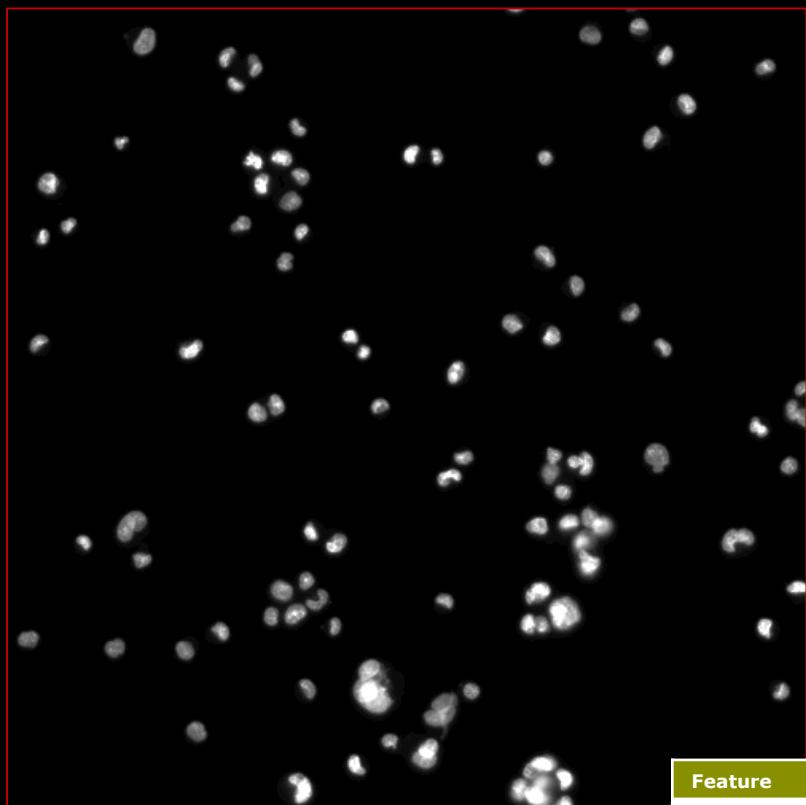
Using the classifier



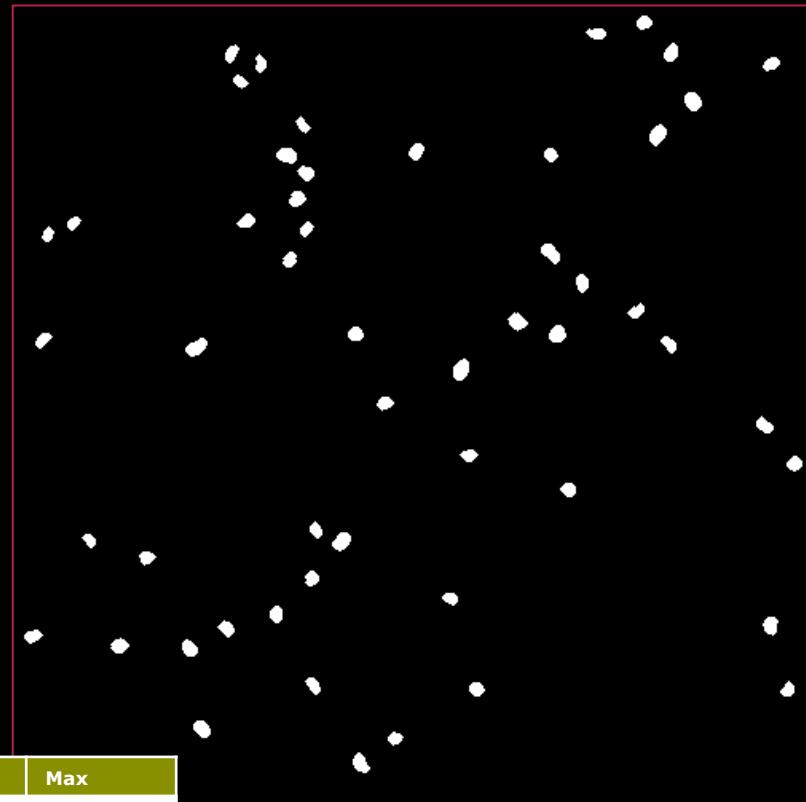
DAPI input image

- Threshold input image
- Morphological opening (SE 5x5)
- Morphological closing (SE 5x5)
- BLOBs found using 8-neighbours
- Border BLOBS removed
- Border features computed
 - Area + circularity
- BLOBs with features inside the acceptance range are **single-nuclei**

Using the classifier



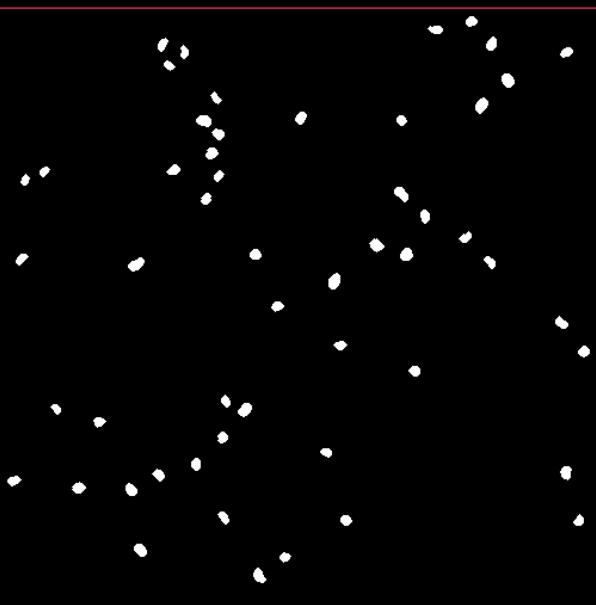
DAPI input image



Found single nuclei

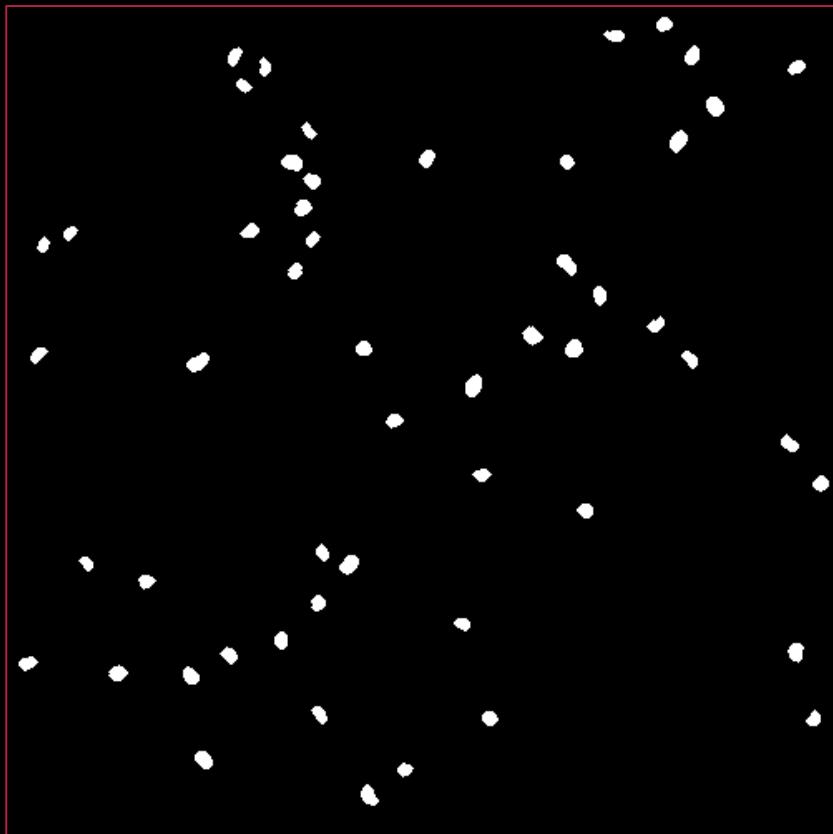
Feature	Min	Max
Area	50	110
Circularity	0.87	1.05

How well does it work?

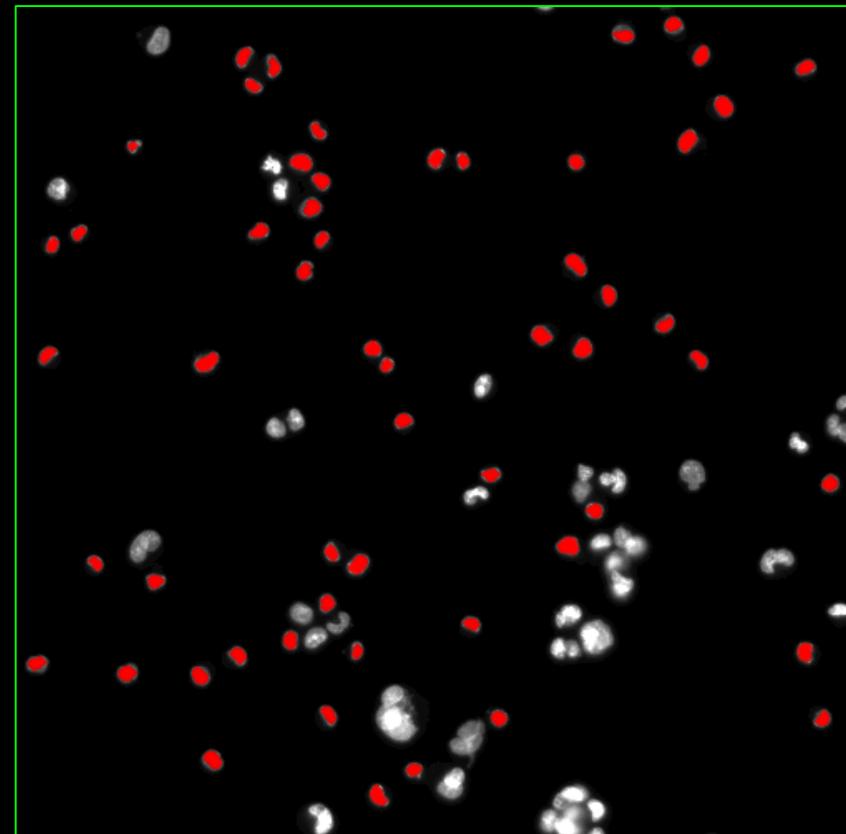


- We say we have a **great** algorithm!
- Strangely the doctor/biochemist do not trust this statement!
 - They need numbers!
- How do we report the performance?

Creating ground truth – expert annotations



Found single nuclei



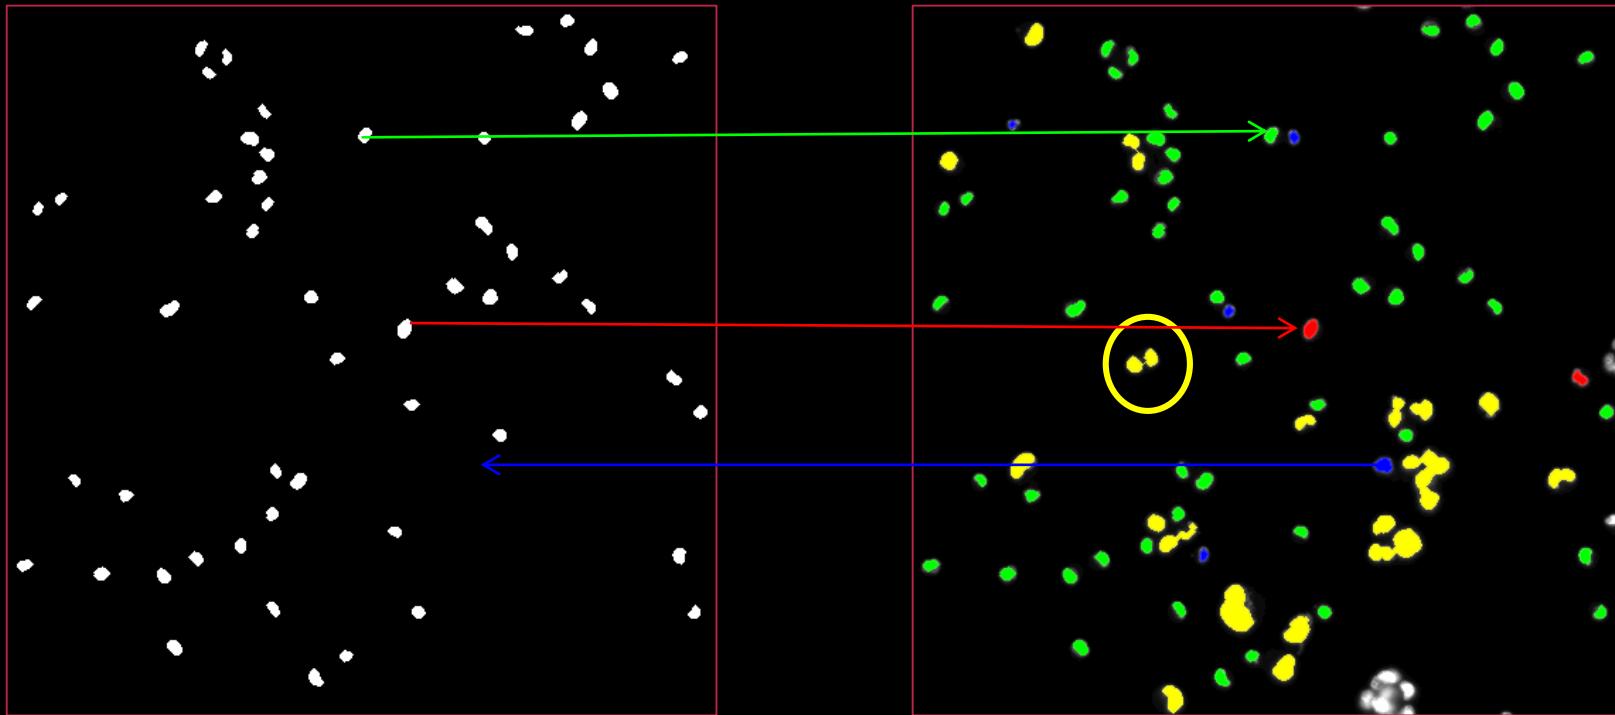
Expert opinion on true single nuclei

Red markings: Single nuclei

Not marked: Noise

Four cases

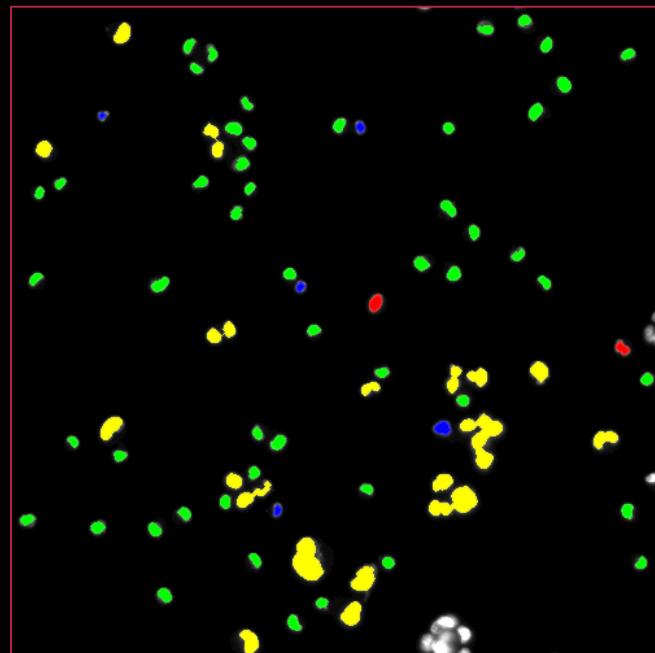
- True Positive (TP): A nuclei is classified as a nuclei
- True Negative (TN): A noise object is classified as noise object
- False Positive (FP): A noise object is classified as a nuclei
- False Negative (FN): A nuclei is classified as a noise object



Found single nuclei

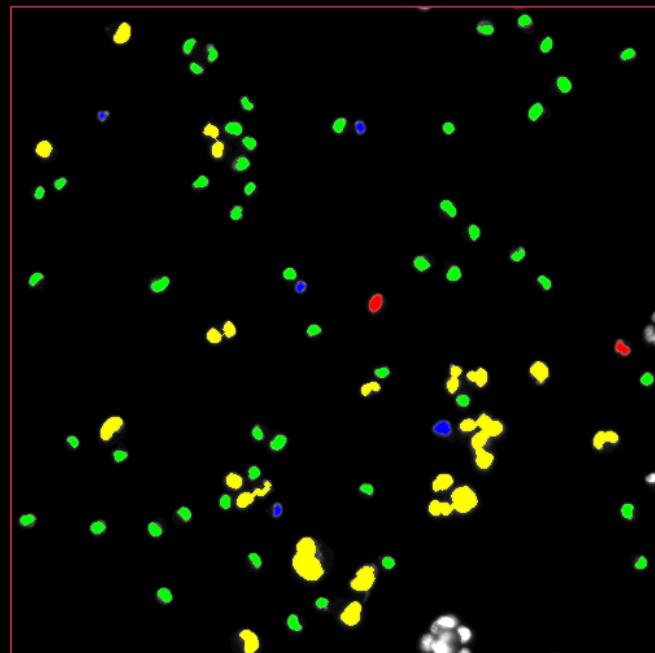
Confusion matrix

	Predicted as noise	Predicted as single-nuclei
Actual noise		
Actual single-nuclei		



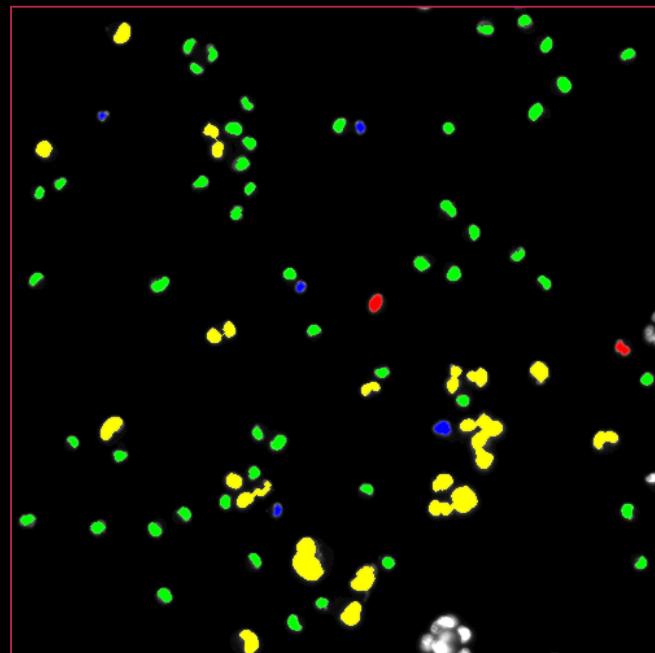
Confusion matrix

	Predicted as noise	Predicted as single-nuclei
Actual noise	TN=19	
Actual single-nuclei		



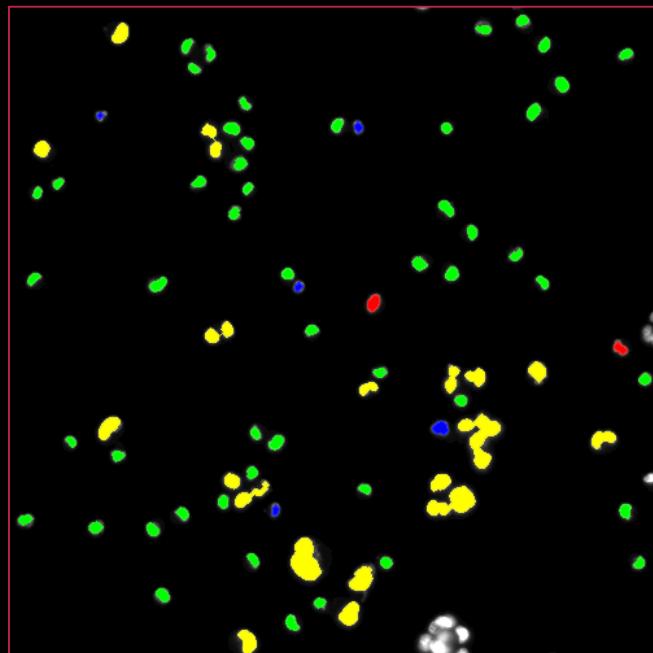
Confusion matrix

	Predicted as noise	Predicted as single-nuclei
Actual noise	TN=19	
Actual single-nuclei		TP=51



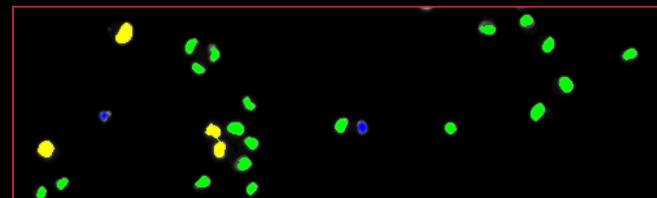
Confusion matrix

	Predicted as noise	Predicted as single-nuclei
Actual noise	TN=19	FP=2
Actual single-nuclei		TP=51

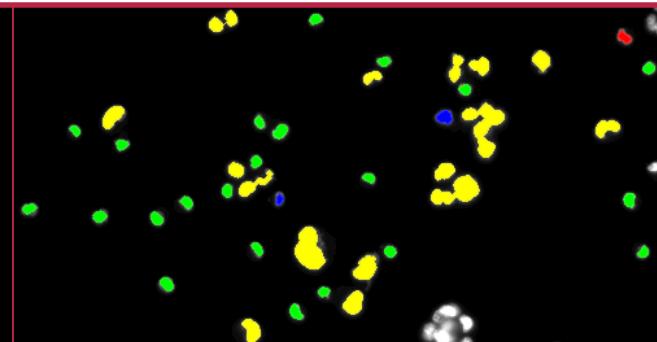


Confusion matrix

	Predicted as noise	Predicted as single-nuclei
Actual noise	TN=19	FP=2
Actual single-nuclei	FN=5	TP=51



Something simpler?



Accuracy

- Tells how often the classifier is correct

$$\text{Accuracy} = \frac{TP + TN}{N}$$

- N is the total number of annotated objects

$$N = TN + TP + FP + FN$$

Accuracy from Confusion Matrix

	Predicted as noise	Predicted as single-nuclei
Actual noise	TN=19	FP=2
Actual single-nuclei	FN=5	TP=51

42%

65%

77%

91%

97%

Accuracy from Confusion Matrix

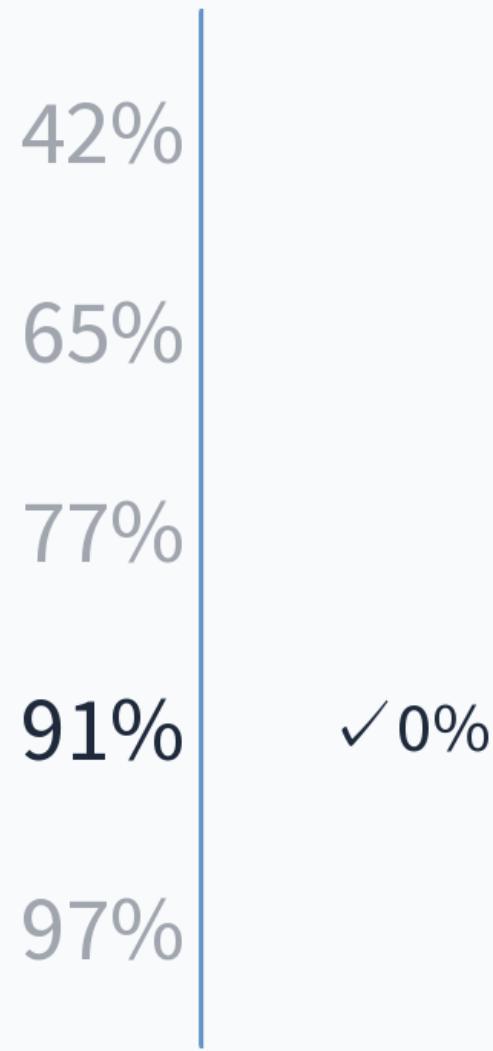
	Predicted as noise	Predicted as single-nuclei
Actual noise	TN=19	FP=2
Actual single-nuclei	FN=5	TP=51

42%
65%
77%
91%
97%

Start the presentation to see live content. For screen share software, share the entire screen. Get help at pollev.com/app

Accuracy from Confusion Matrix

	Predicted as noise	Predicted as single-nuclei
Actual noise	TN=19	FP=2
Actual single-nuclei	FN=5	TP=51



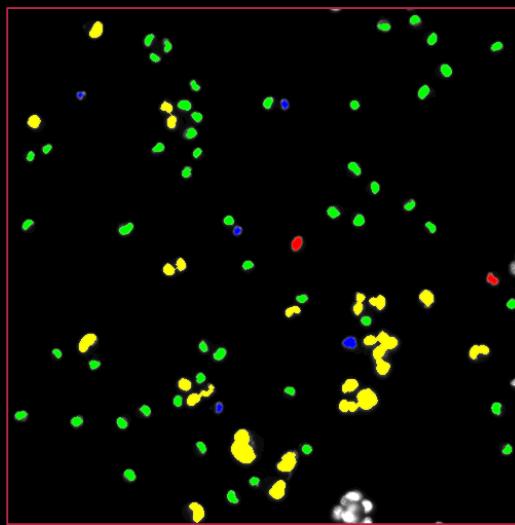
Start the presentation to see live content. For screen share software, share the entire screen. Get help at pollev.com/app

True positive rate (sensitivity)

- How often is a positive predicted when it actually is positive

$$\text{Sensitivity} = \frac{TP}{FN+TP}$$

All the experts true single-nuclei



Sensitivity from Confusion Matrix

	Predicted as noise	Predicted as single-nuclei
Actual noise	TN=19	FP=2
Actual single-nuclei	FN=5	TP=51

62%

65%

71%

91%

93%

Start the presentation to see live content. For screen share software, share the entire screen. Get help at pollev.com/app

Sensitivity from Confusion Matrix

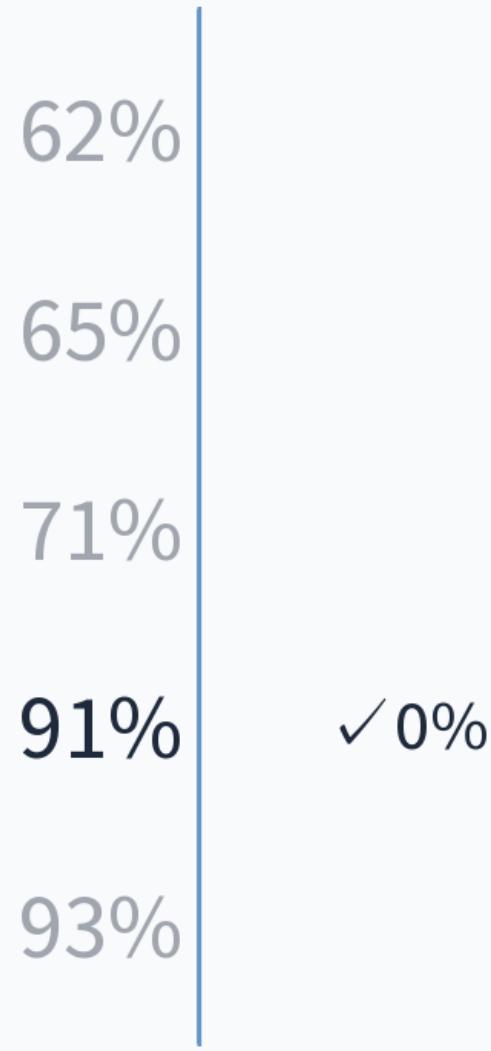
	Predicted as noise	Predicted as single-nuclei
Actual noise	TN=19	FP=2
Actual single-nuclei	FN=5	TP=51

62%
65%
71%
91%
93%

Start the presentation to see live content. For screen share software, share the entire screen. Get help at pollev.com/app

Sensitivity from Confusion Matrix

	Predicted as noise	Predicted as single-nuclei
Actual noise	TN=19	FP=2
Actual single-nuclei	FN=5	TP=51



Start the presentation to see live content. For screen share software, share the entire screen. Get help at pollev.com/app

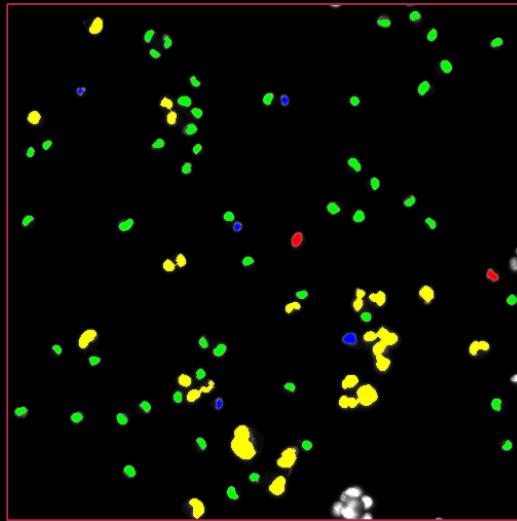
Specificity

- How often is a negative predicted when it actually is negative

$$\text{Specificity} = \frac{TN}{TN + FP}$$

TN + FP

All the experts true noise objects



True positive rate

You have made an algorithm that can locate neon fish in an aquarium. An expert has marked all neon fish in an image as seen in Figure 1 (left). The result of your algorithm is seen in Figure 1 (right). What is the true positive rate of your algorithm?

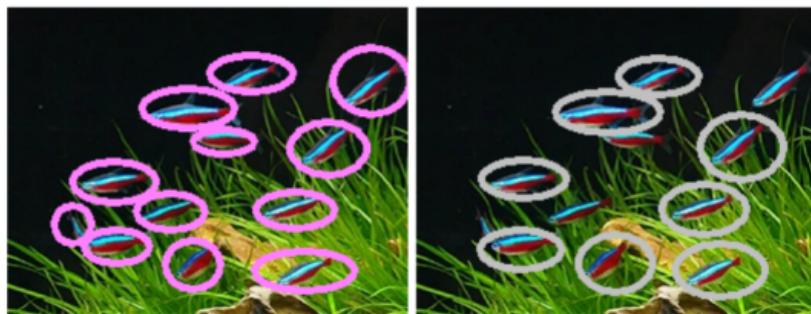


Figure 1: Image of aquarium with neon fish. Left: Expert markings are shown as ellipses. Right: Algorithm markings are shown as ellipses.

77%

92%

81%

55%

67%

True positive rate

You have made an algorithm that can locate neon fish in an aquarium. An expert has marked all neon fish in an image as seen in Figure 1 (left). The result of your algorithm is seen in Figure 1 (right). What is the true positive rate of your algorithm?

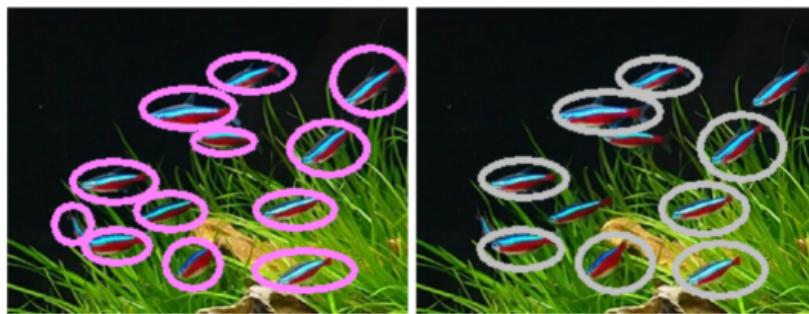
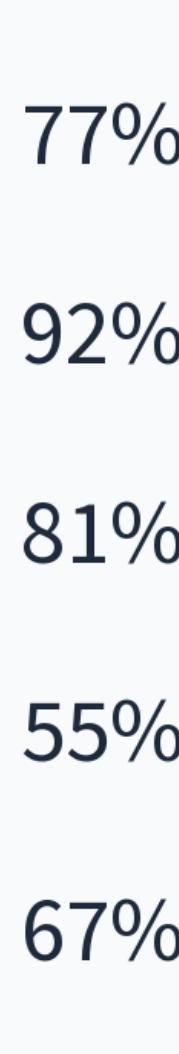


Figure 1: Image of aquarium with neon fish. Left: Expert markings are shown as ellipses. Right: Algorithm markings are shown as ellipses.



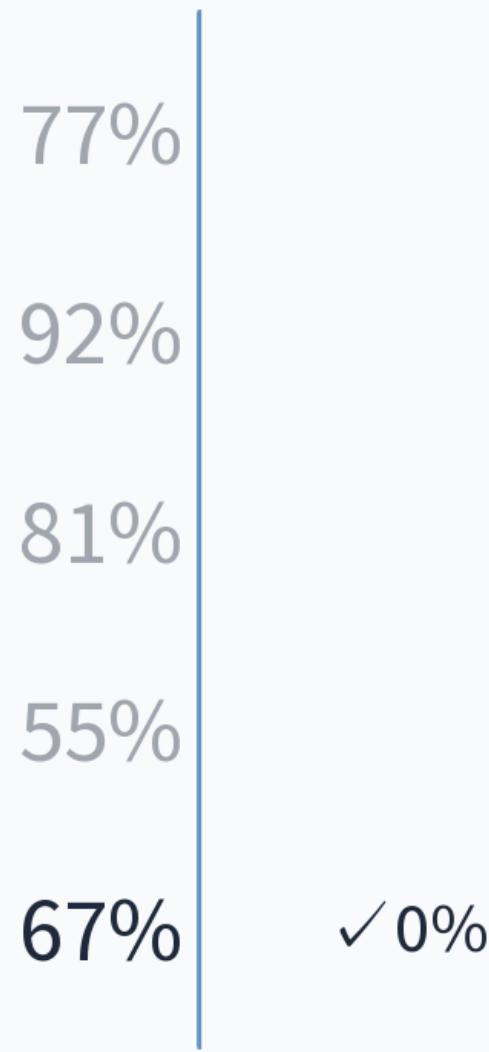
Start the presentation to see live content. For screen share software, share the entire screen. Get help at pollev.com/app

True positive rate

You have made an algorithm that can locate neon fish in an aquarium. An expert has marked all neon fish in an image as seen in Figure 1 (left). The result of your algorithm is seen in Figure 1 (right). What is the true positive rate of your algorithm?

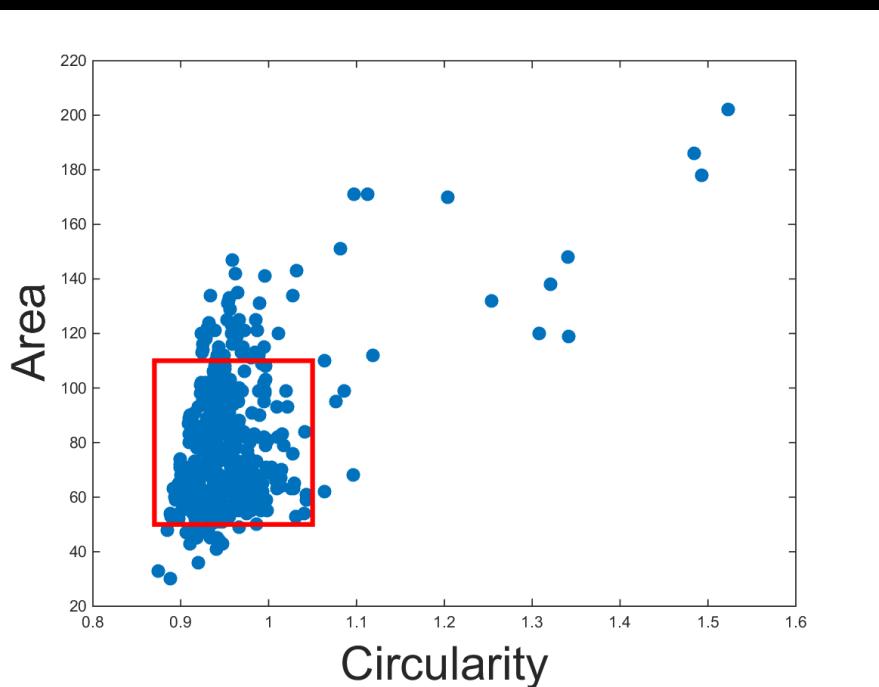


Figure 1: Image of aquarium with neon fish. Left: Expert markings are shown as ellipses. Right: Algorithm markings are shown as ellipses.



Start the presentation to see live content. For screen share software, share the entire screen. Get help at pollev.com/app

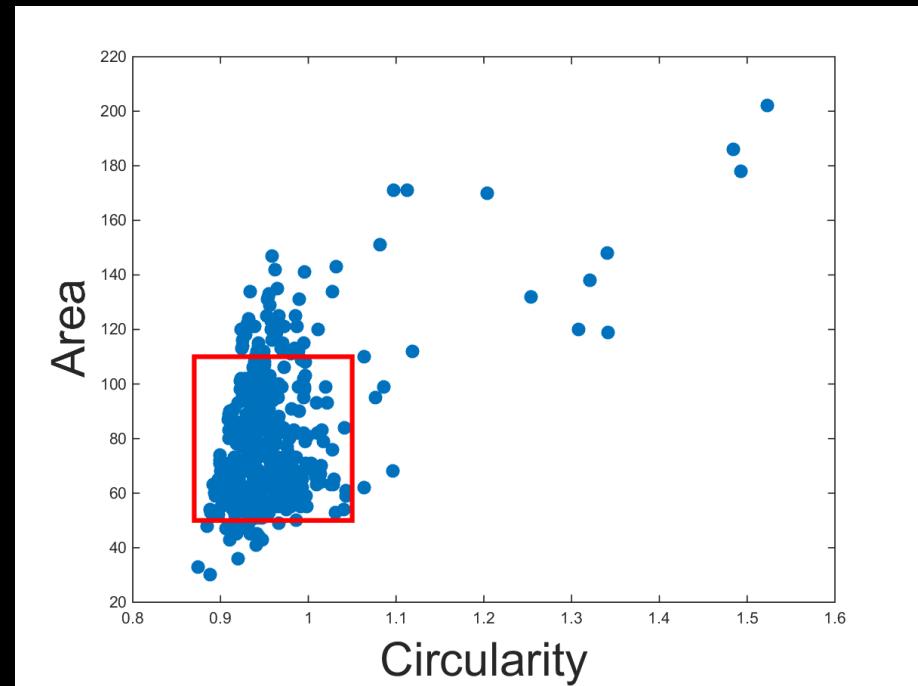
Optimising the classification



- Changing the classification limits
- The rates will be changed:
 - Accuracy
 - Sensitivity
 - Specificity
 - ...
- Very dependent on the task what is optimal

Dependencies

- Increasing **true positive rate**
 - Increased **false positive rate**
 - Decreased **precision**



Example – cell analysis

- We want **only** single-nuclei cells
 - For further analysis
- We **do not** want to do an analysis of a noise object
- We are **not** interested in the true number of single nuclei

What measure is the most important?

- We want **only** single-nuclei cells
 - For further analysis
- We **do not** want to do an analysis of noise objects
- We are **not** interested in the true number of single nuclei

Low false positives

High true positives

High true negatives

Low false negatives

What measure is the most important?

- We want **only** single-nuclei cells
 - For further analysis
- We **do not** want to do an analysis of noise objects
- We are **not** interested in the true number of single nuclei

Low false positives

High true positives

High true negatives

Low false negatives

What measure is the most important?

- We want **only** single-nuclei cells
 - For further analysis
- We **do not** want to do an analysis of noise objects
- We are **not** interested in the true number of single nuclei

Low false positives

✓ 0%

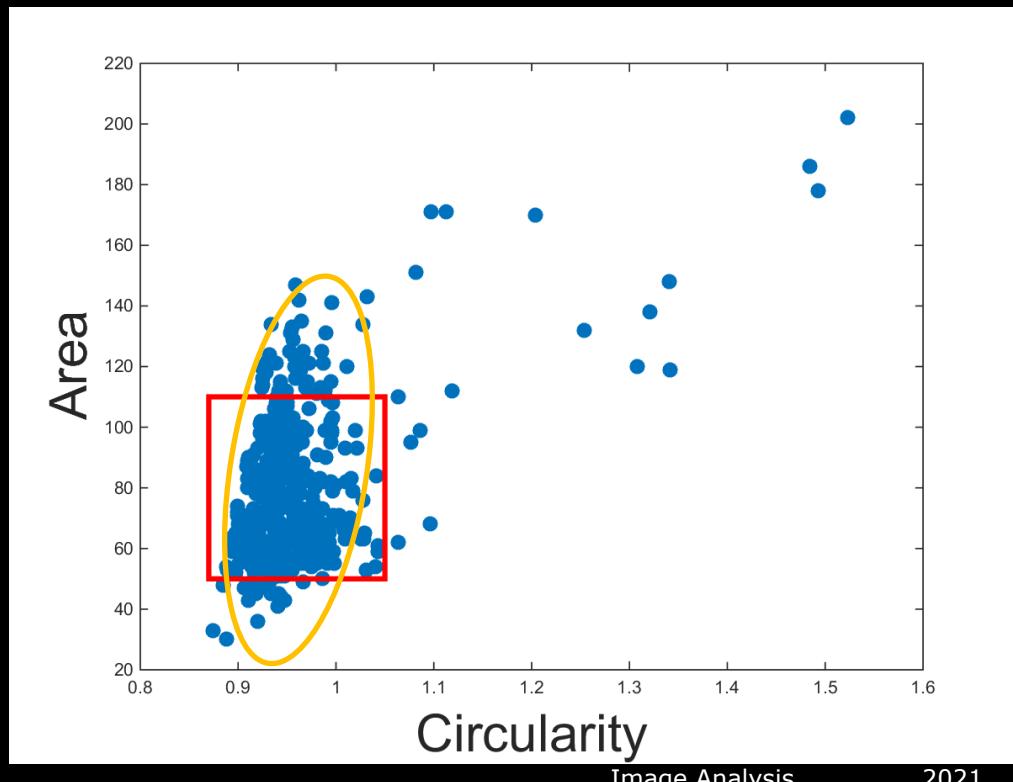
High true positives

High true negatives

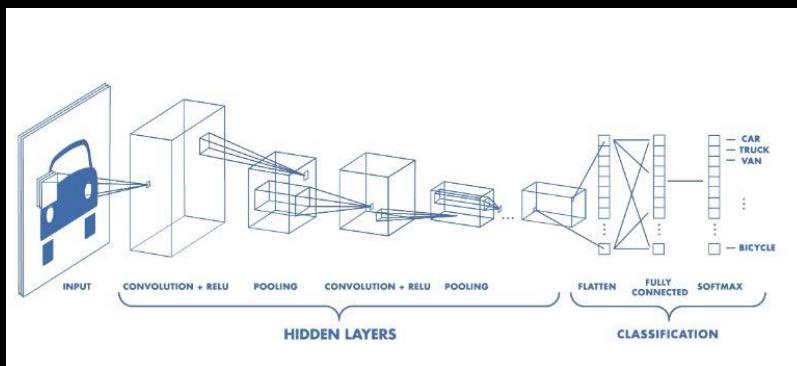
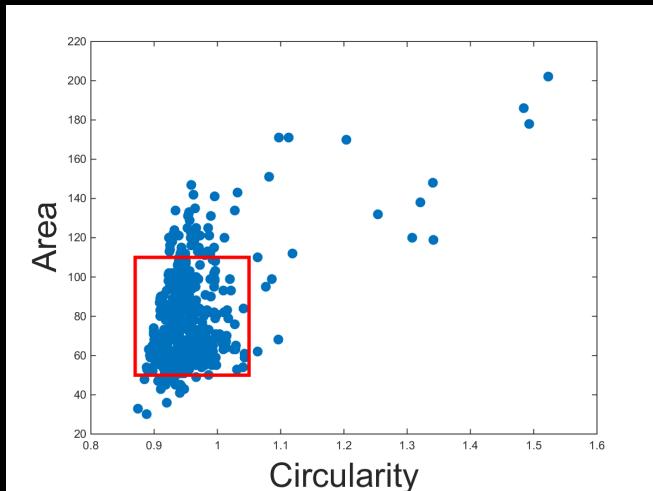
Low false negatives

Advanced classification

- Fitting more advanced functions to the samples
- Multivariate Gaussians
- Mahalanobis distances

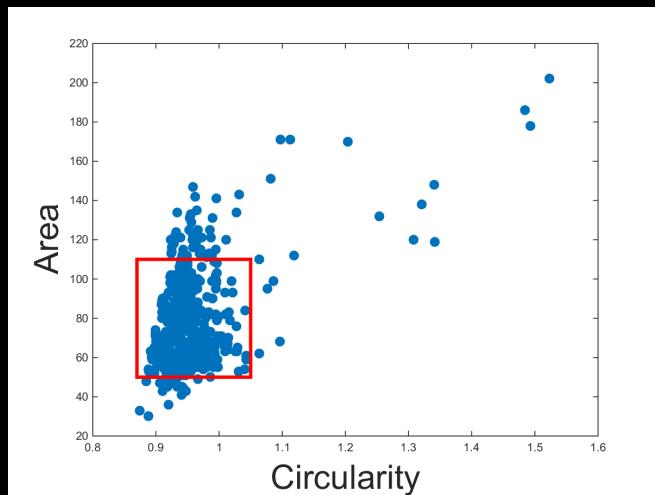


Feature Engineering vs. Deep learning



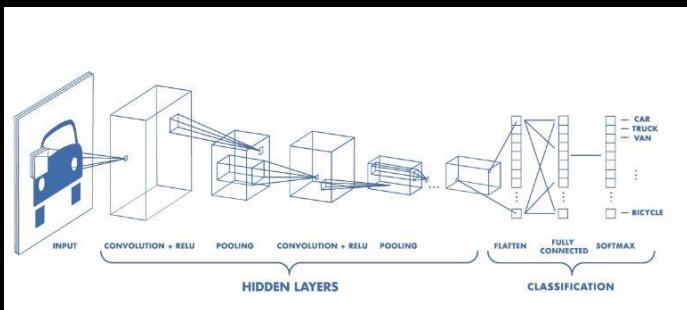
- Until around 5-7 years ago **feature engineering** was the way to go
- Now **deep learning beats** everything
- However – feature engineering is still important

Feature engineering



- Given a classification problem
 - Cars vs. Pedestrians
- Use background knowledge to select relevant features
 - Area
 - Shape
 - Appearance
 - ...
- Use multivariate statistics to classify
- Depending on the selected features

Deep learning

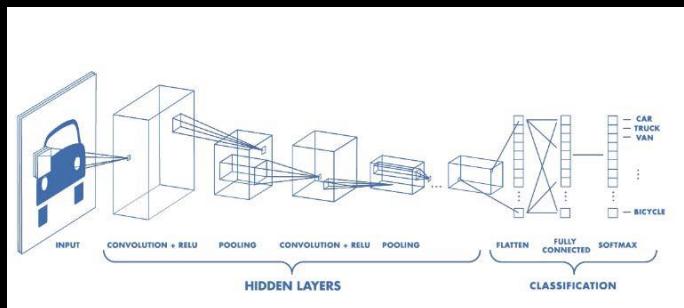


- You start with a dummy classifier
- Feed it with lots and lots of data with given labels
- The network learns the optimal features
- Layer/network engineering

Feature Engineering vs. Deep learning

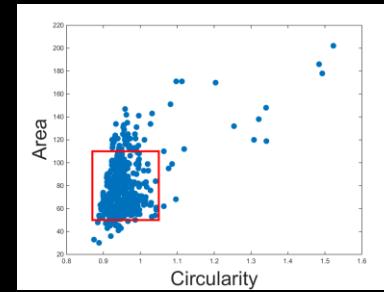
Deep Learning

- When you have lot of annotated data
- Where it is not clear what features work



Manual features

- When you have limited data
- When it is rather obvious what features can discriminate



The level of the lecture

Far too easy - my hamster
could understand it

Too easy - I need more

Suitable - I am generally
learning what I want

Too hard - slow down
please

Far too hard - my head is
exploding

The quizzes

Not enough quizzes -
I want more more

Fine with the quizzes
- no more no less

Argghh! These
quizzes...I want less

Next week

- Geometric transformations
- Landmark based registration